Create Your Own Intelligent Voice Assistant

Presented by Petr Grinberg
HSE, 2025
Based on HSE DLA Course: https://github.com/markovka17/dla

Welcome! I am Petr

- ***MSc EPFL, LauzHack Mentor***
- ***Lecturer at NRU HSE***
- ***Ex. Reality Defender Inc.***
- ***Ex. Samsung R&D***

Follow me on GitHub (@Blinorot)

Plan of the workshop

1. Deep Learning for Audio Processing
    a. A bit of history
    b. What can we do with it?
2. Signal Processing Basics
3. The core parts of a Voice Assistant
    a. KWS
    b. ASR
    c. TTS
    d. LLM
4. The code for all parts:
    a. Vanilla PyTorch
    b. Pre-trained models
    c. APIs

Slides

A short history of Speech Recognition

50's

In **1952**, *Bell Laboratories* designed the "**Audrey**" system which could recognize a single voice speaking **digits** aloud

60's

In **1962,** *IBM* introduced "**Shoebox**" which understood and responded to **16 words** in English.

A short history of Speech Recognition

70's

DARPA's system was capable of understanding over **1,000** words. **Siri** was a spin-out of DARPA development :)

80's

The '80s saw speech recognition vocabulary go from a few hundred words to **several thousand words** thanks to **HMM**

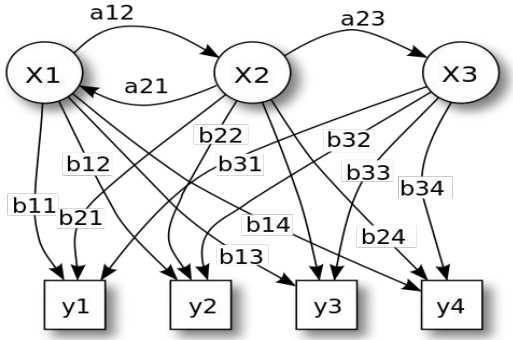

A short history of Speech Recognition

90's

Speech recognition was propelled forward in the 90s in large part because of **faster processors**

00-10's

And then came the era of big data, machine learning and GPUs

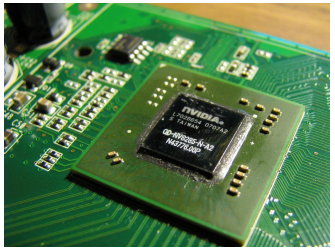

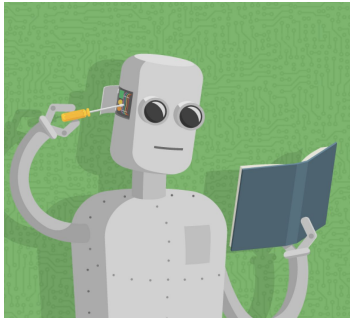

A short history of Speech Synthesis

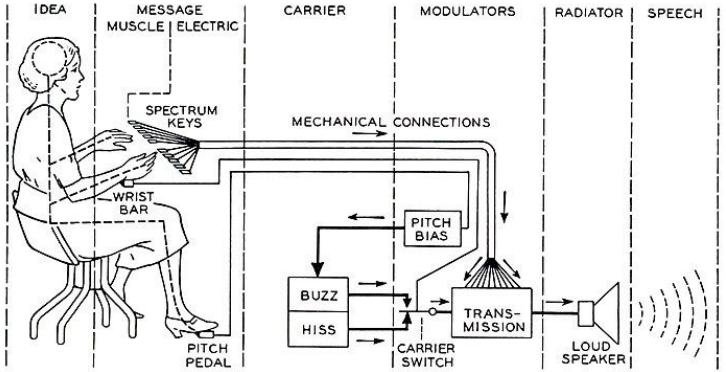

Fig. 8—Schematic circuit of the voder.

30's

In **1939**, *The Bell Laboratory's* **Voder** was the first attempt to electronically synthesize human speech by breaking it down into its **acoustic components**

until 80's

Formant-based on rules. You may listen examples in Atari&Sega games :)

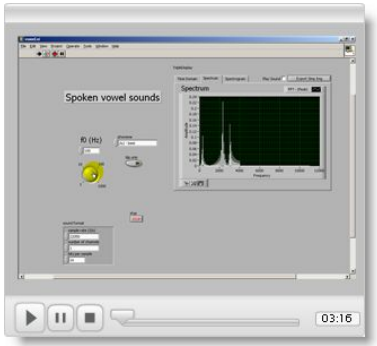

A short history of Speech Synthesis

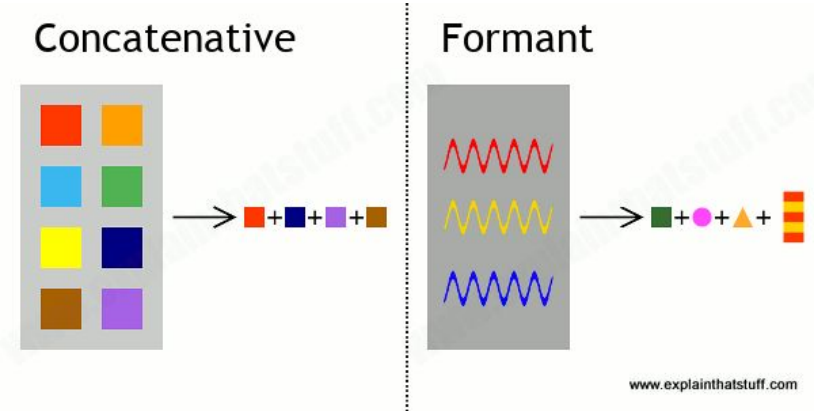

90–00's

Concatenative synthesis is a technique for synthesising sounds by concatenating short samples of recorded sound (called *units*).

10's

And then came the era of big data, machine learning and GPUs

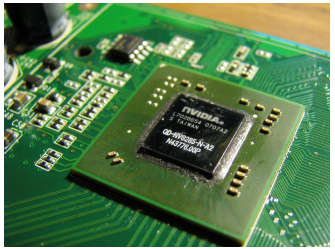

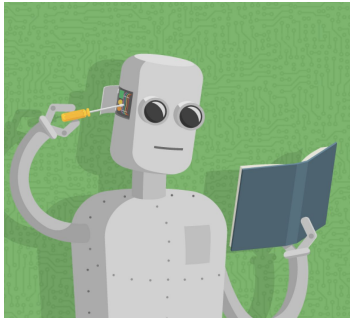

Tasks: Keyword Spotting (KWS)

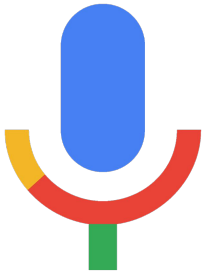

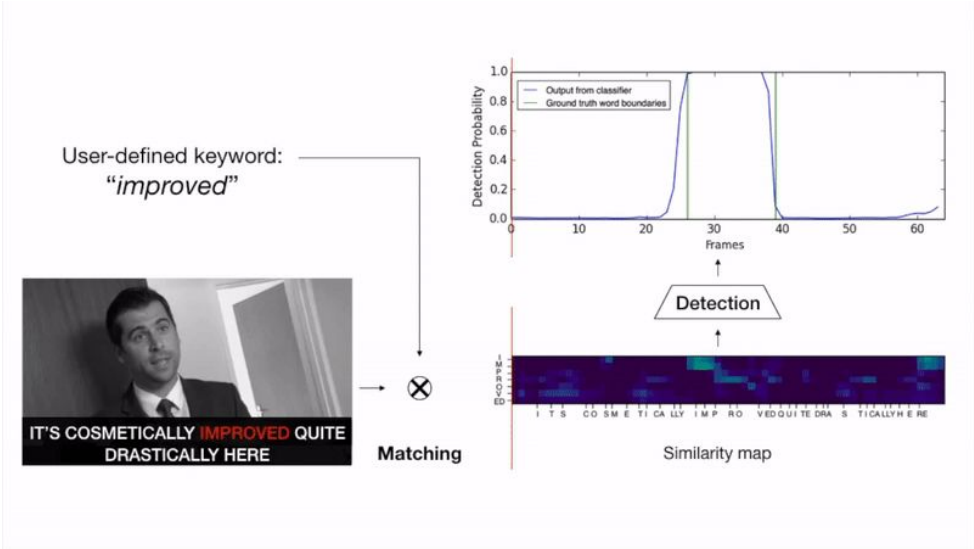

Momeni, Liliane, et al. "Seeing wake words: Audio-visual keyword spotting." arXiv preprint arXiv:2009.01225 (2020).

Tasks: Automatic Speech Recognition

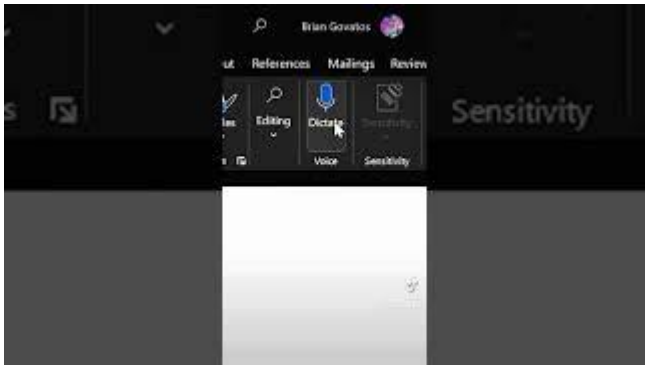

Voice Dictation

Meeting Summary

Speech to Text

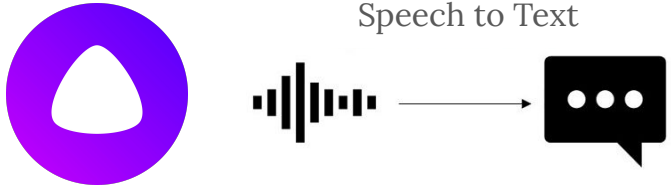

And, of course, voice assistants

Tasks: Text To Speech And Voice Conversion

Audio (Audio-Visual) Deepfakes*

AI Covers

Text to Speech

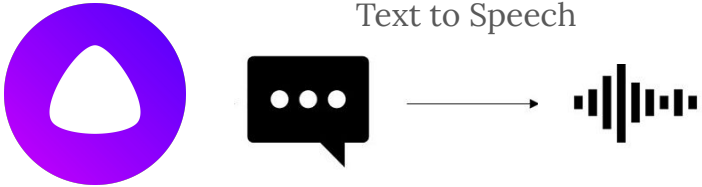

And, of course, voice assistants

*SadTalker + XTTS

Tasks: Video To Audio

https://deepmind.google/discover/blog/generating-audio-for-video/

Voice Technologies Mind Map

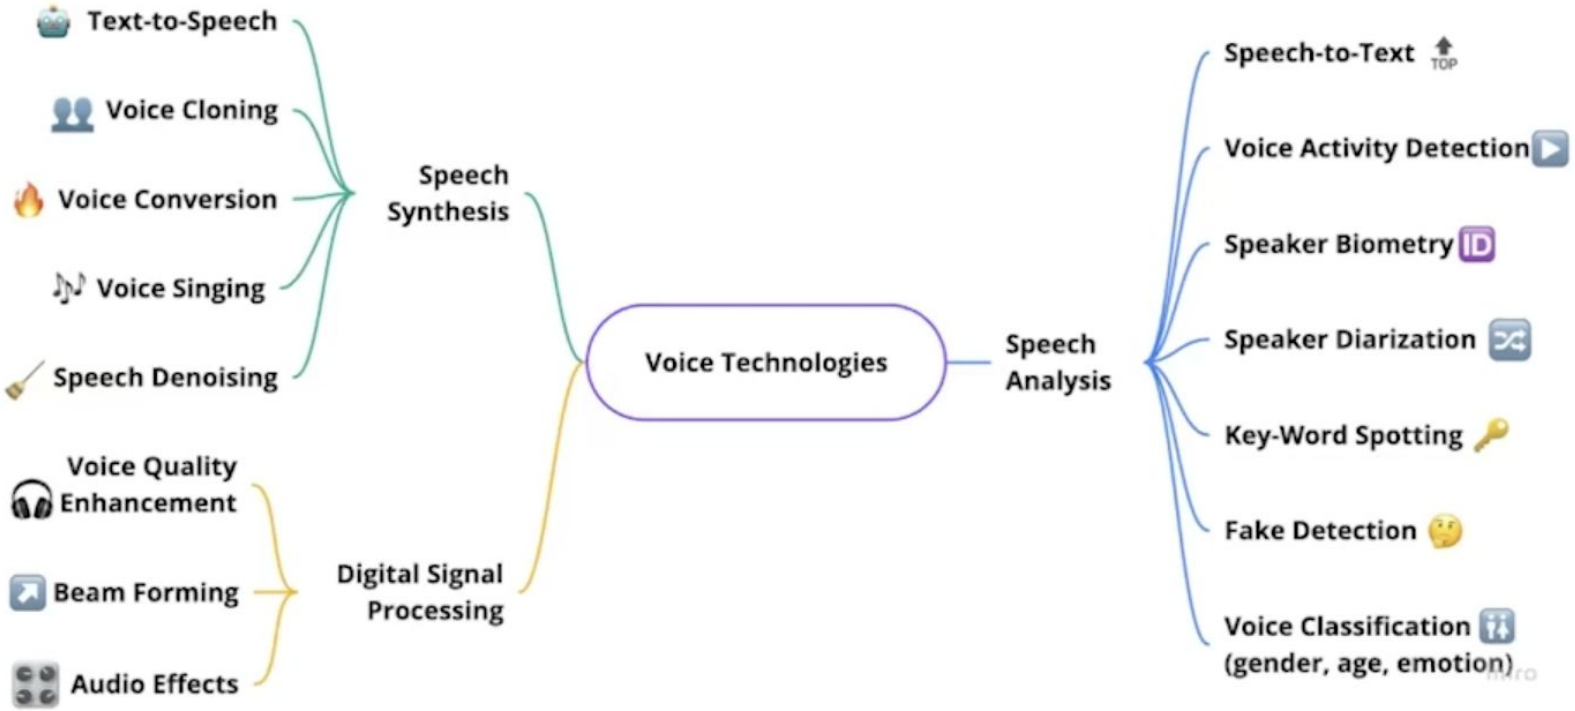

https://github.com/yandexdataschool/speech_course

Voice Technologies Mind Map

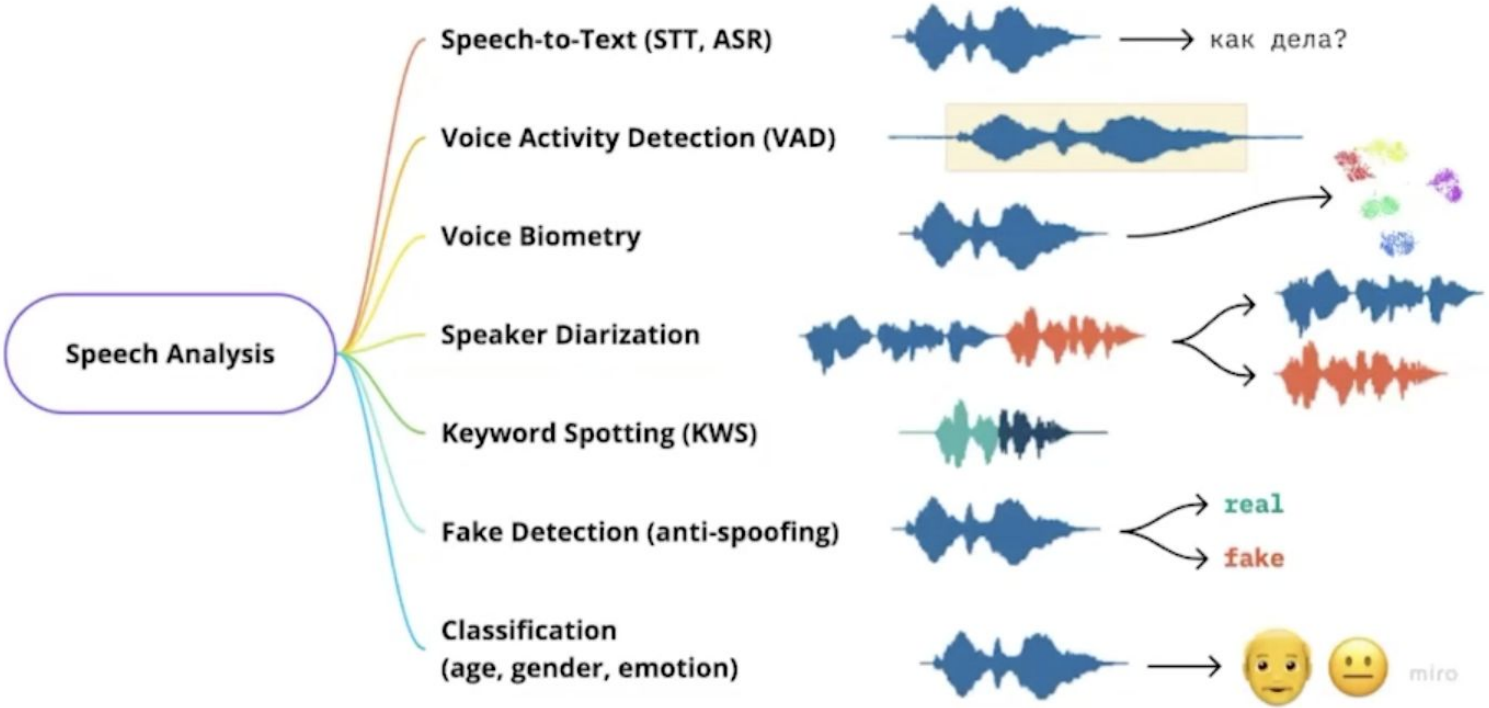

https://github.com/yandexdataschool/speech_course

Voice Technologies Mind Map

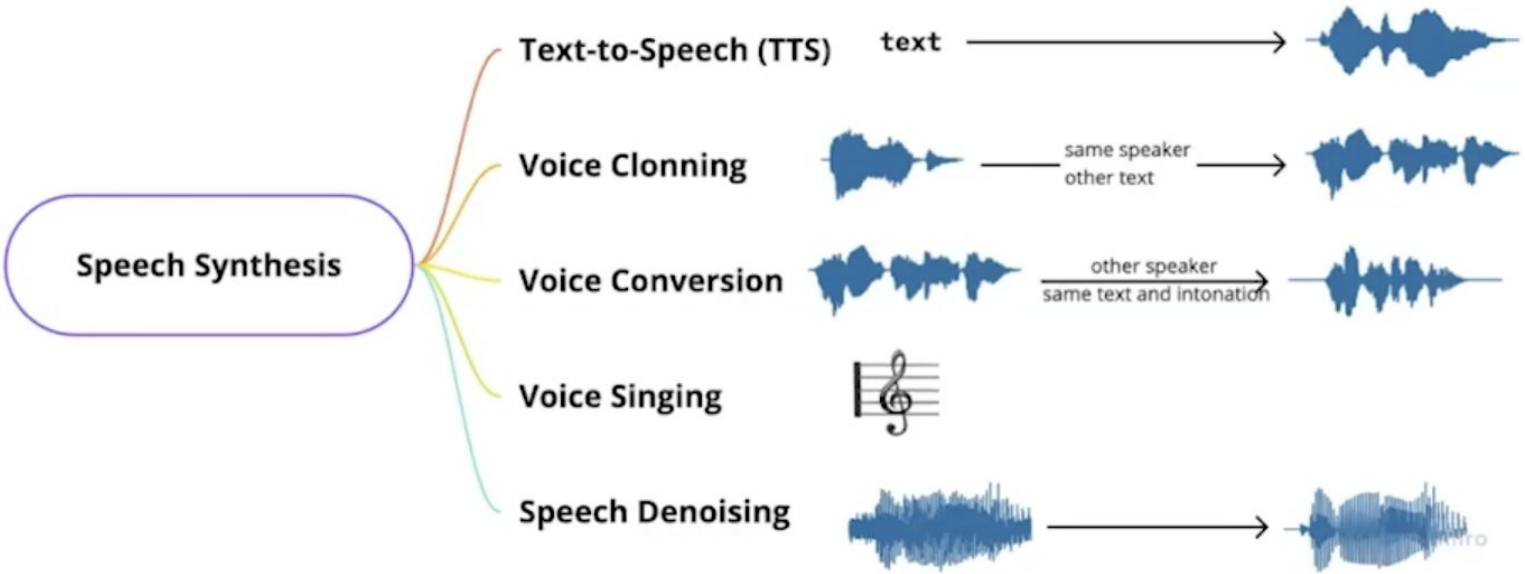

https://github.com/yandexdataschool/speech_course

Brief Background for Core Assistant Parts

What is sound?

- **Sound wave** is the pattern of **oscillations** caused by the movement of energy traveling through the air
- **Microphone** picks up these air **oscillations** and converts them into electrical vibrations
- These **oscillations** are converted into an **analog** signal and then into a **digital** signal

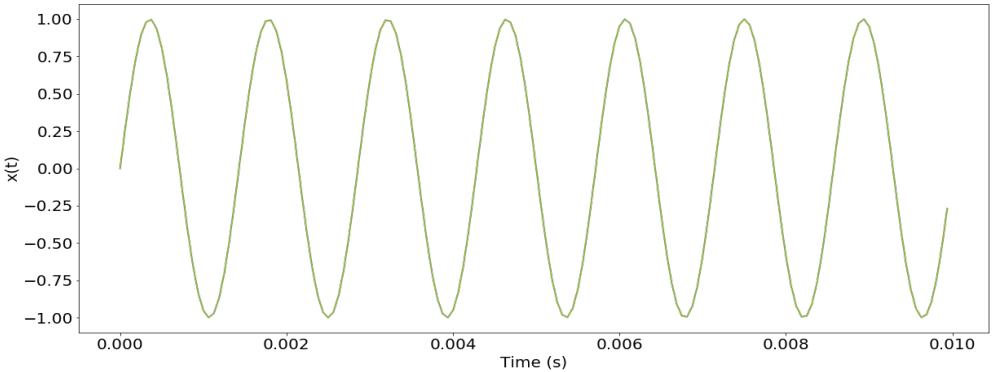

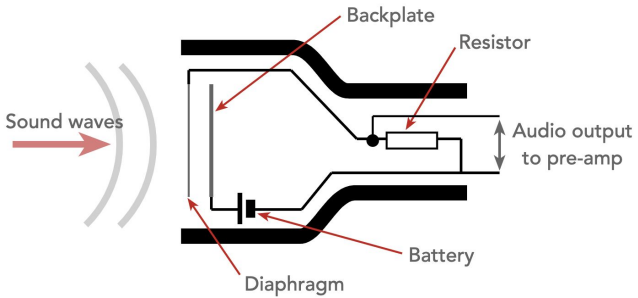

Analog-to-Digital Conversion

- Converting analog signals to a sequence of numbers having finite precision
- Corresponding devices are called A/D converters (ADCs)

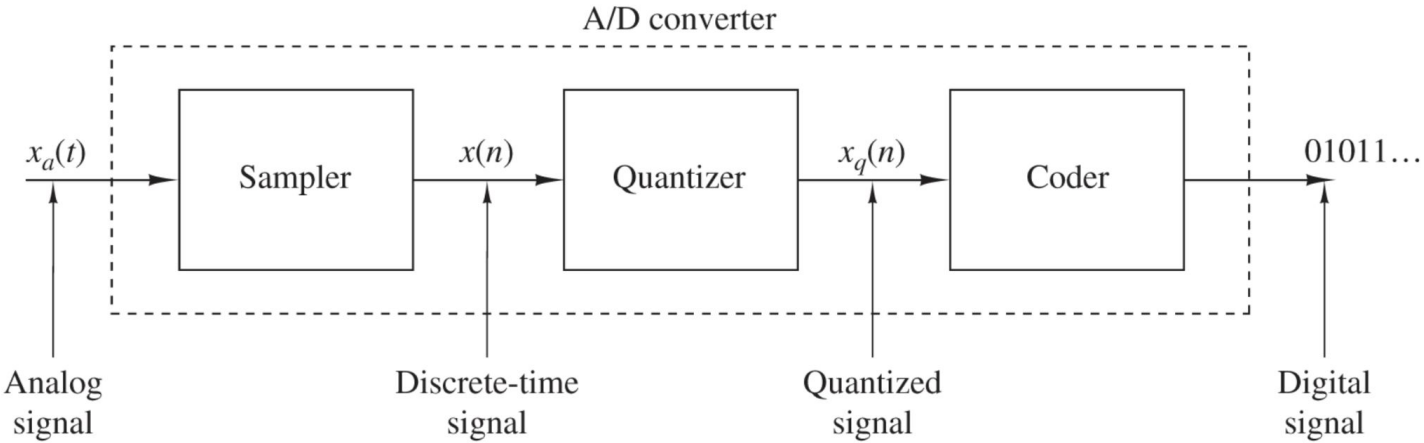

Analog-to-Digital Conversion

- X[n] = X(nT), where T is the **sampling period** in seconds
- **Sampling rate** (Hz) f = 1 / T, the number of audio samples in 1 sec.

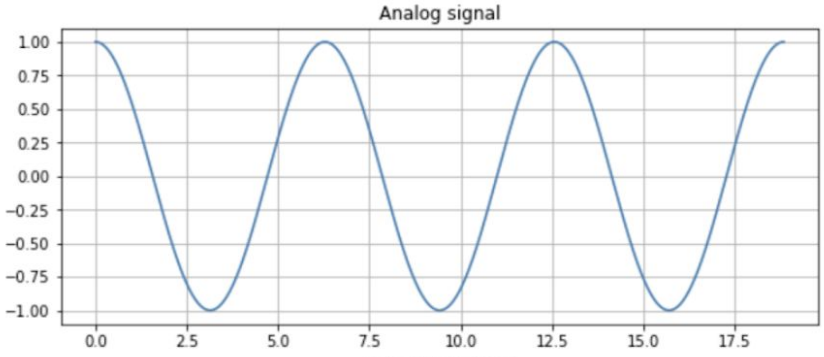

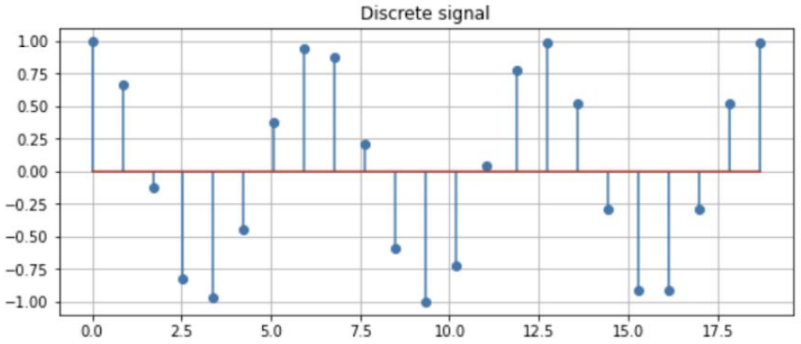

Usually Models are trained on recordings with a specific Sample rate. You should remember it and Resample audio before feeding it to the network.

Analog-to-Digital Conversion

- Using some interpolation techniques, we can convert digital signal back to analog one

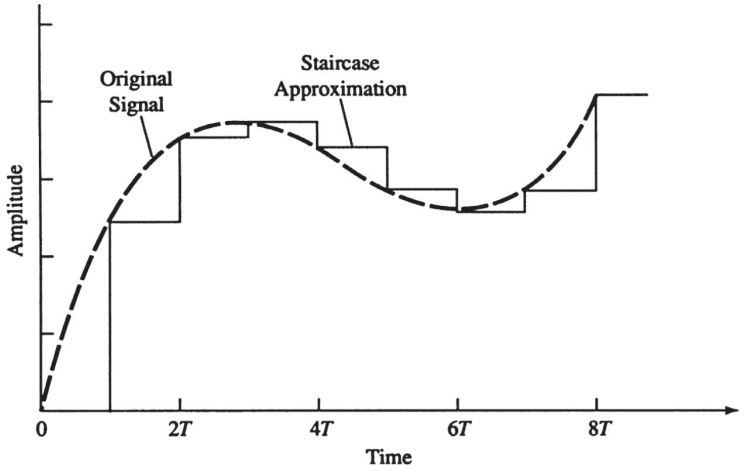

Spectrograms

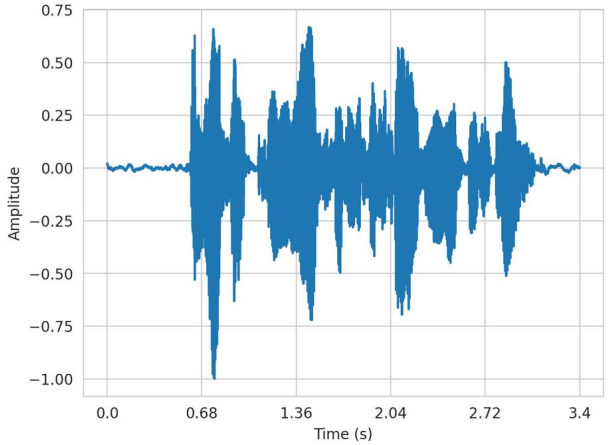

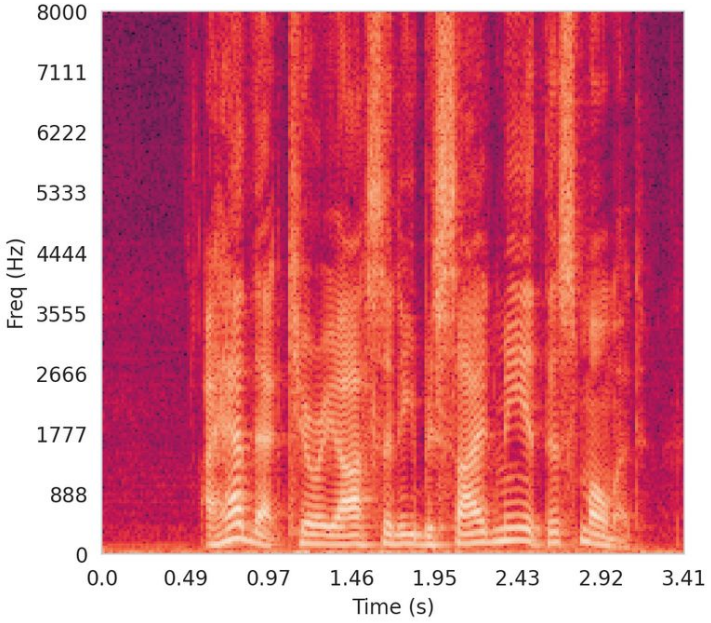

Problems with Raw Audio (Waveform) representation

1 second of audio contains Sample Rate elements in it. This leads to severe computational costs

- For example, 4 minutes of audio with 44.1kHz sample rate results in a 10M-length sequence.

- In contrast, high resolution 1024x1024 image will require only 3M elements (in RGB)

Problems with Raw Audio (Waveform) representation

One letter (sound) requires thousands of samples (example for 22050 Hz sample rate):

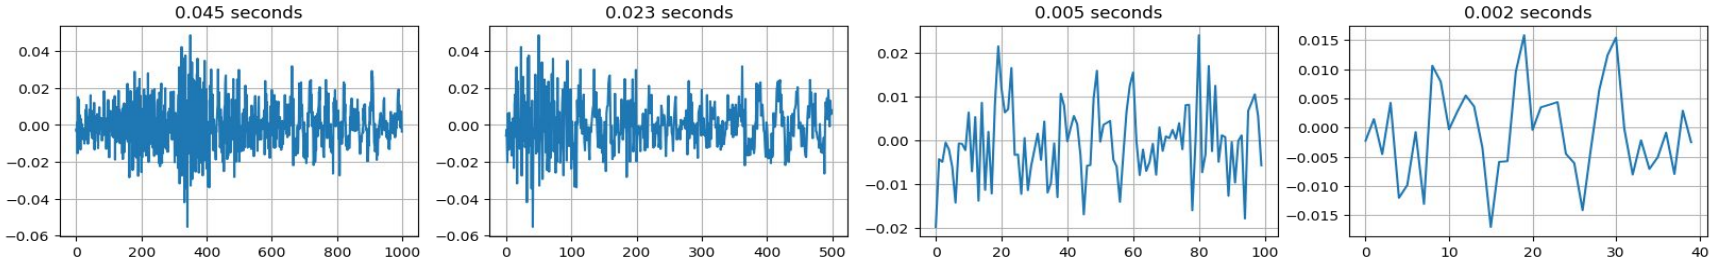

Visualization is similar after transforms (like noise or pitch shift, etc) – hard to say anything about transform from the visualization

Voice

Voice + Noise

Frequencies

We need something apart from **temporal** representation to better understand what happens inside the audio, how it was modified, etc.

Frequency domain – another domain to represent audio whether some things may become clearer

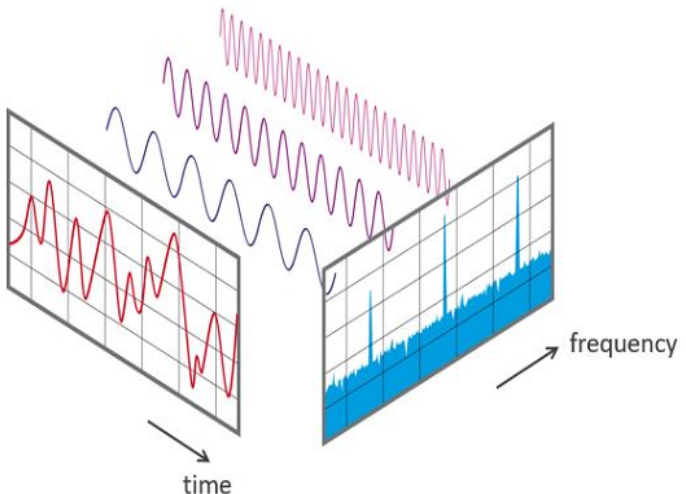

Frequencies – Fourier Transform

Sounds can be modeled as sum of sinusoid
→ determine frequencies and amplitudes

To find these sinusoids, we apply **Fourier Transform**, which is basically a **change of basis** in the vector space.

Fourier transform integral

$$\hat{f}(\xi) = \int_{-\infty}^{\infty} f(x)\, e^{-i2\pi\xi x}\, dx, \quad \forall\, \xi \in \mathbb{R}.$$

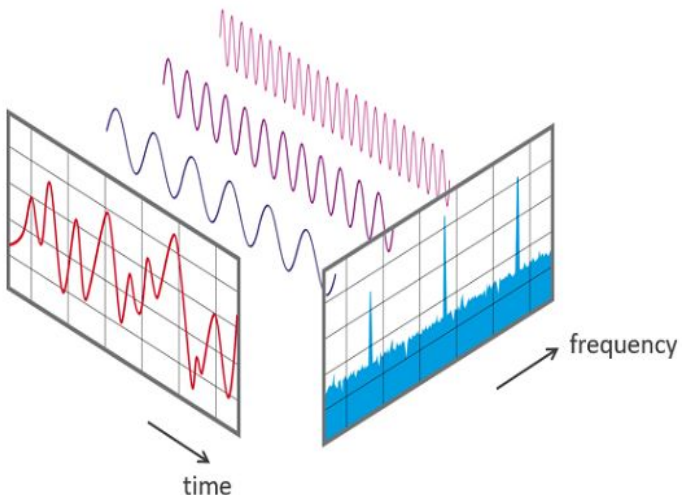

Frequencies – Fourier Transform

Let's look at the visualization.

Remark 1: the output of fourier transform is complex (not real) by design.

Remark 2: In audio DL, we usually use discrete version

$$X_k = \sum_{n=0}^{N-1} x_n \cdot e^{-i 2\pi \frac{k}{N} n}$$

Frequencies – Fourier Transform

Note that Spectrum is symmetric for real inputs

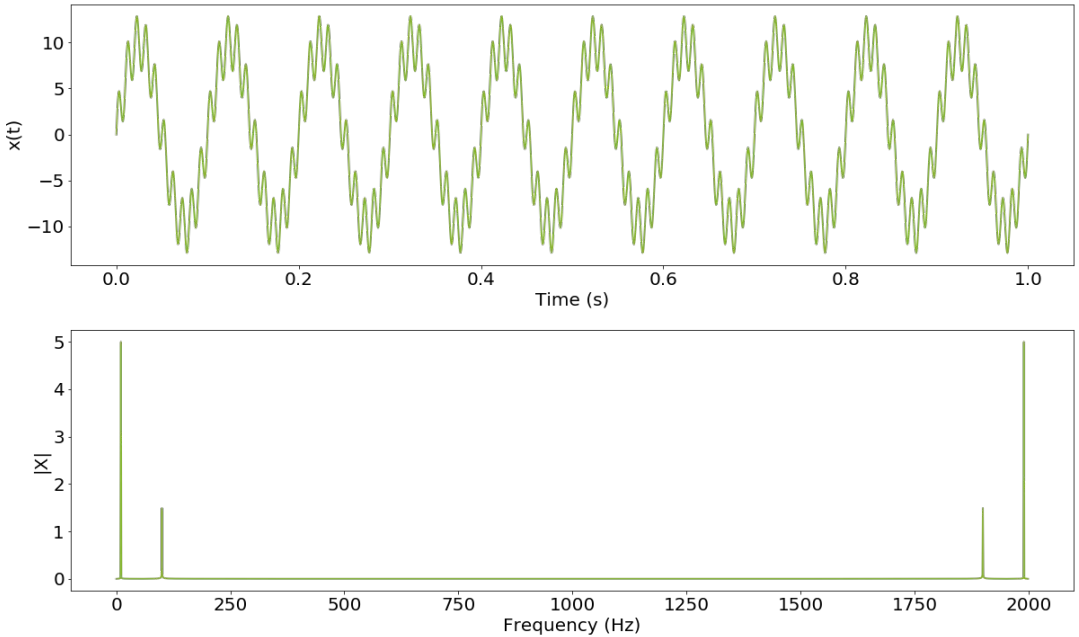

$$F = 2kHz$$

$$f(t) = 10\sin(2\pi 10t) + 3\sin(2\pi 100t)$$

Frequencies – Fourier Transform

That's why we usually work with only half of it

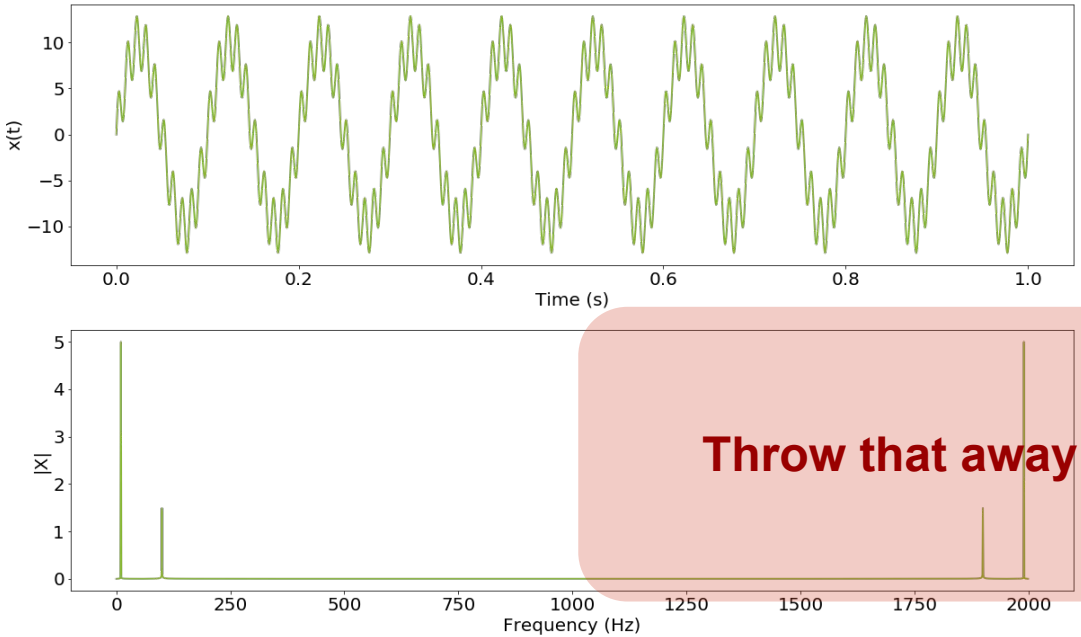

$$F = 2kHz$$

$$f(t) = 10\sin(2\pi 10t) + 3\sin(2\pi 100t)$$

Frequencies – Fourier Transform

Some stuff, like noise, may be easier to distinguish in the frequency domain:

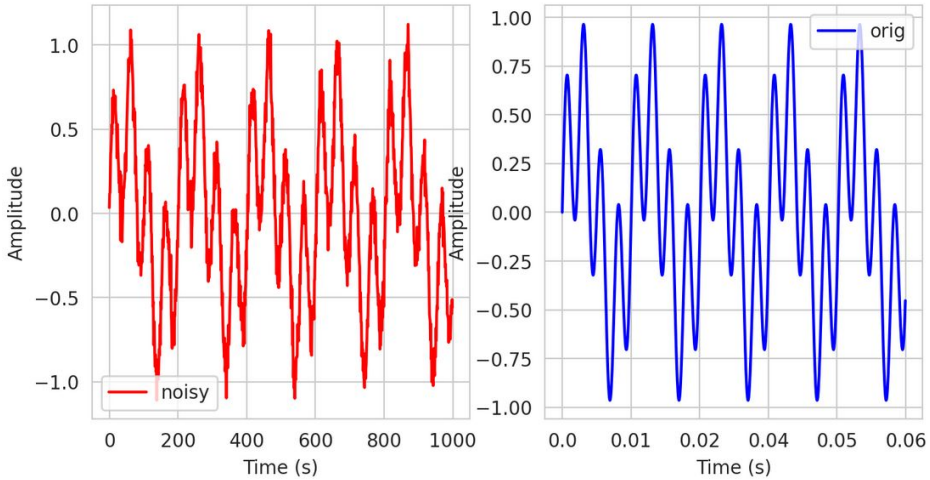

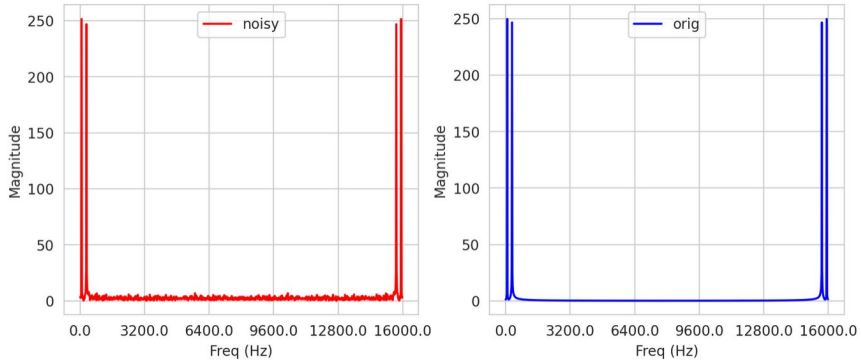

Time-Frequency Domain

Frequencies may rapidly change through time. So we want to see these changes

That's why, in practice, we work in Time-Frequency domain, where we see both changes across time and changes across frequencies

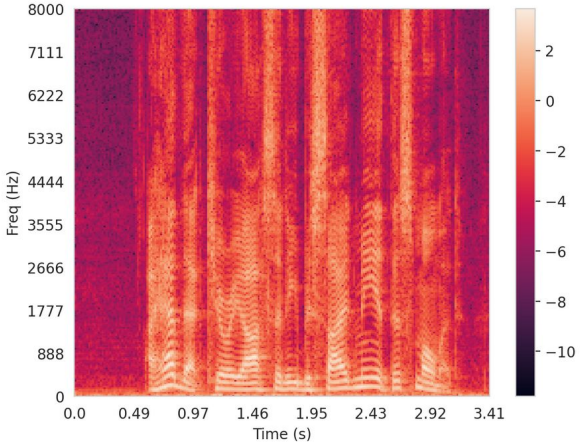

Short-Time Fourier Transform

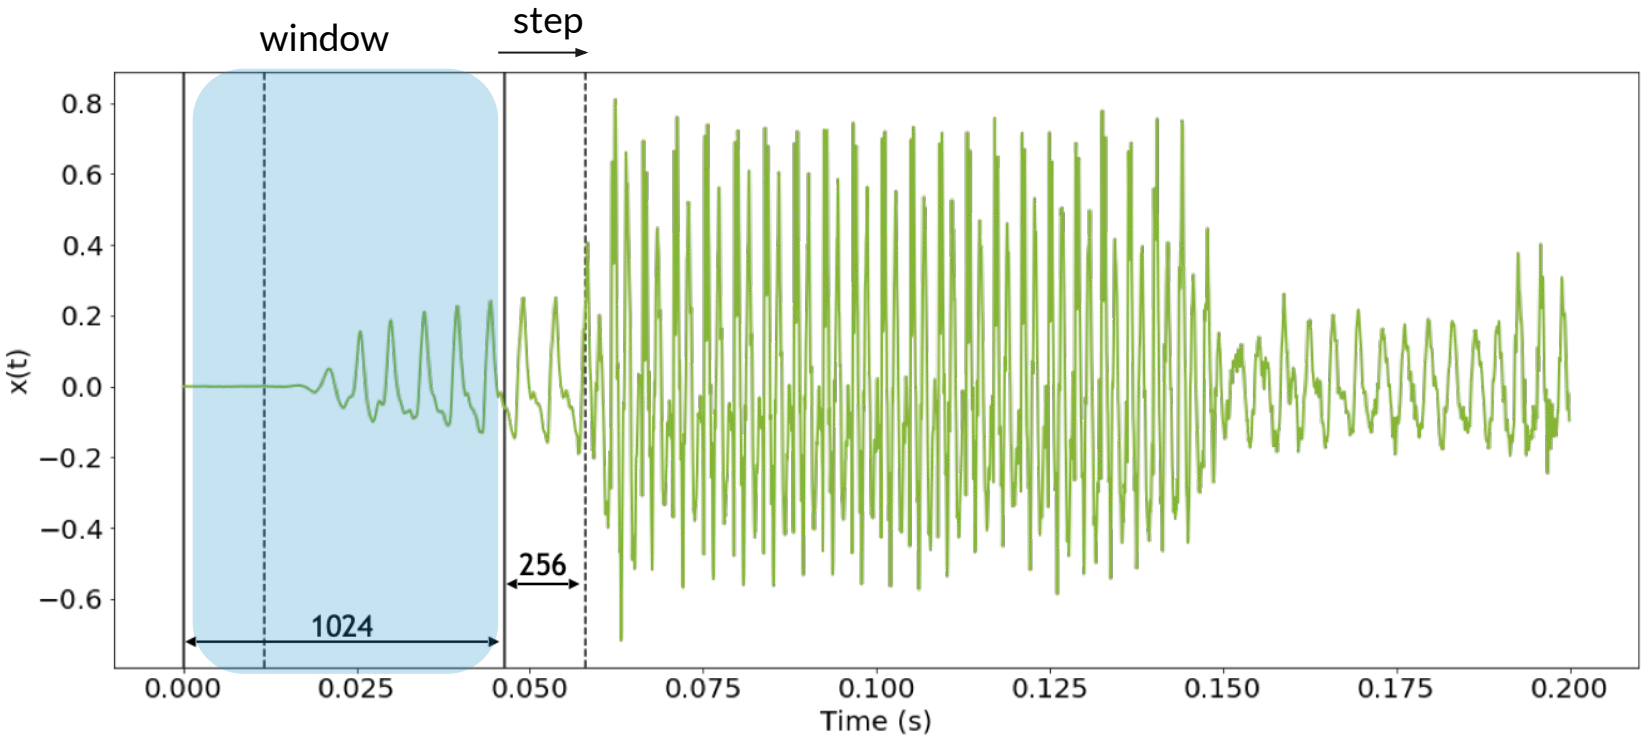

Short-Time Fourier Transform

Due to overlapping, we use window function to smooth the boundaries

Sliced signal

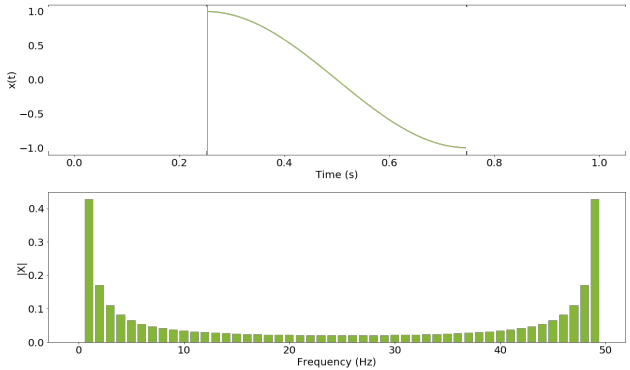

Window

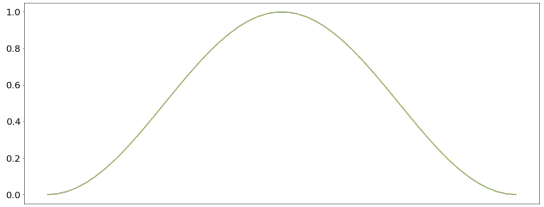

Windowed signal

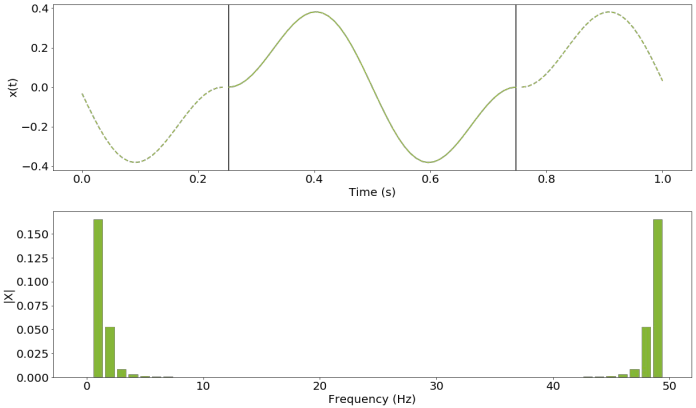

Short-Time Fourier Transform

So we apply DFT for each window and stack the outputs

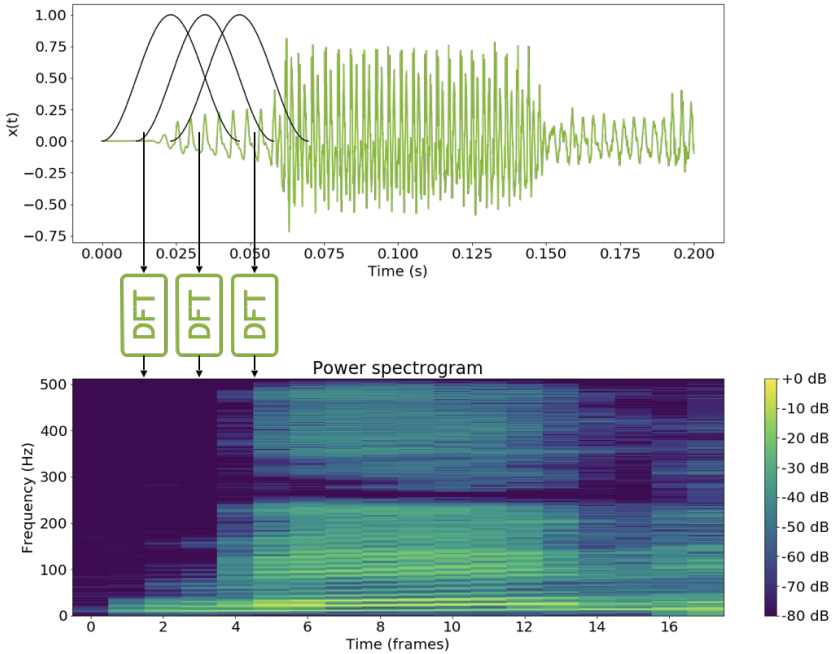

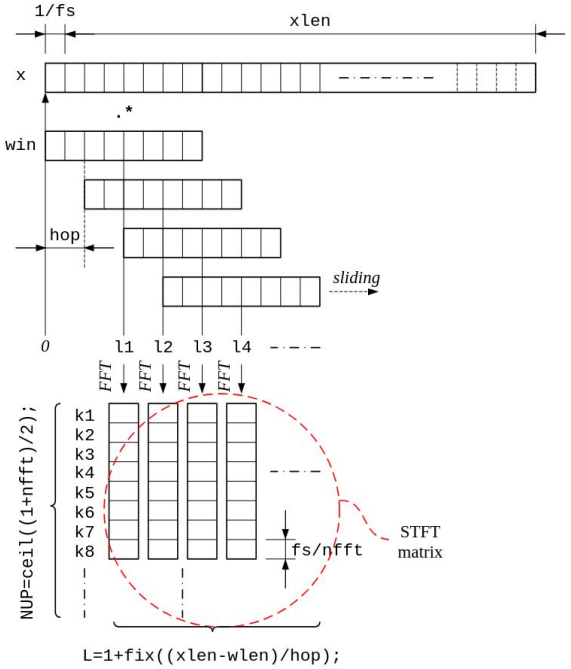

Spectrogram

Notice that the output of STFT is **complex-valued**. Usually we look at magnitude and phase of the output.

The magnitude version is called **Magnitude Spectrogram or just Spectrogram**. The phase version is called **Phase Spectrogram**.

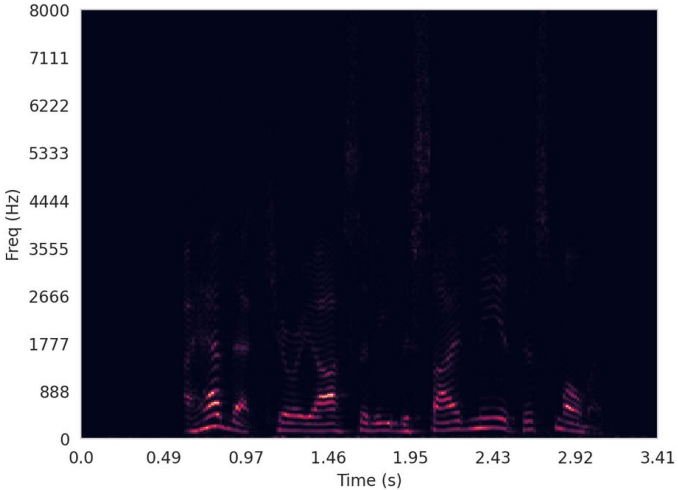

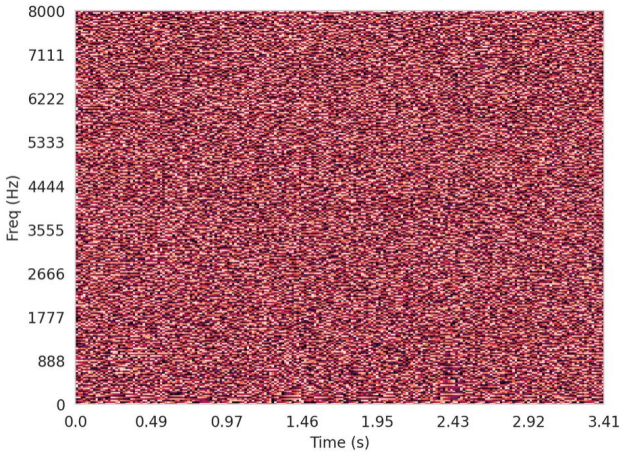

Spectrogram

Also, we often calculate the logarithm of the spectrogram, because it is much easier to interpret

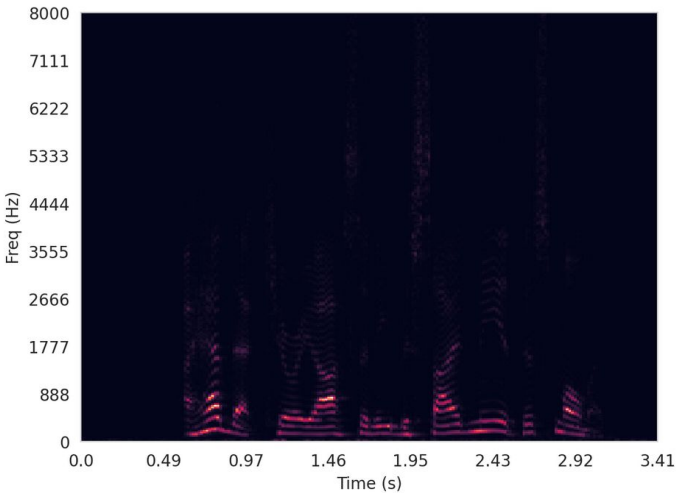

log

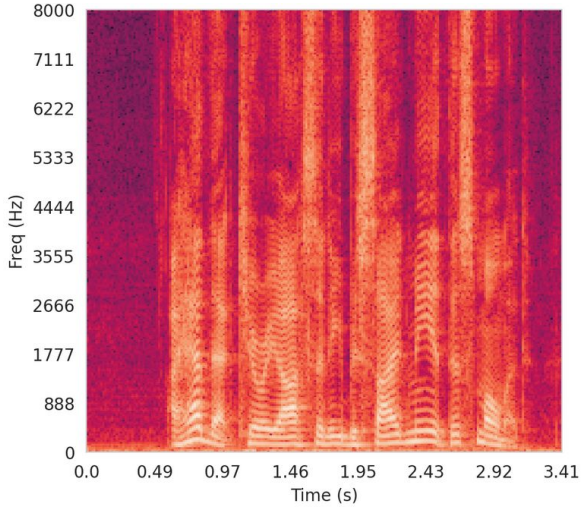

Short-Time Fourier Transform

STFT also allows us to see noise, etc. The speech signal, for example, will contain harmonics which can be clearly seen in the STFT

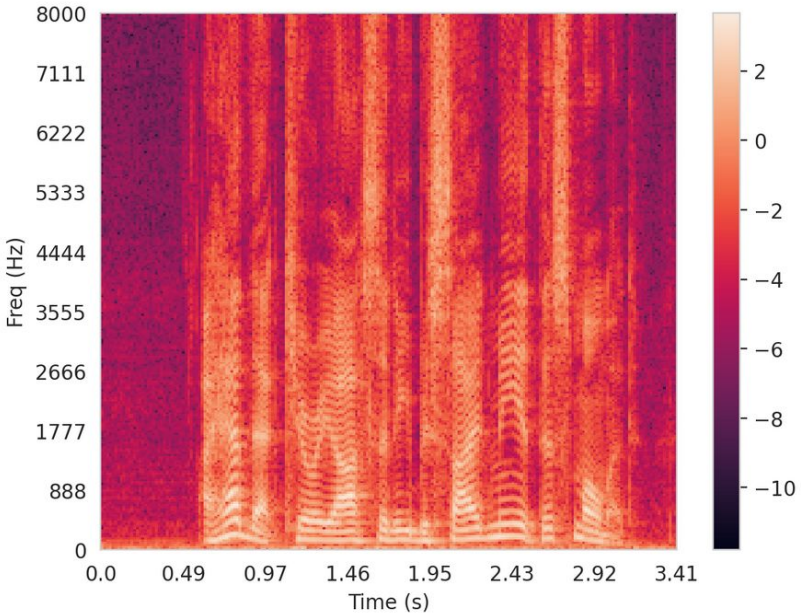

By looking at harmonics it is possible to detect what letter was said in this period!

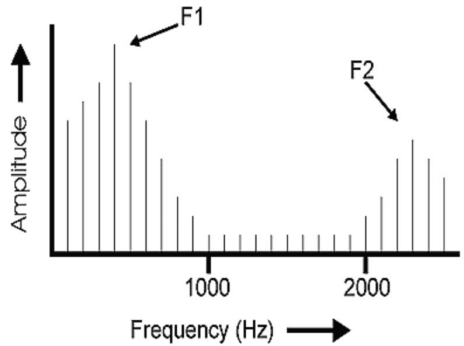

Remark: Inverse Transform

Fourier Transform is a change of basis. So it is invertible transform.

Given some constraints on window and hop (step) size, the STFT transform will also be **lossless** and we can invert to get original raw audio.

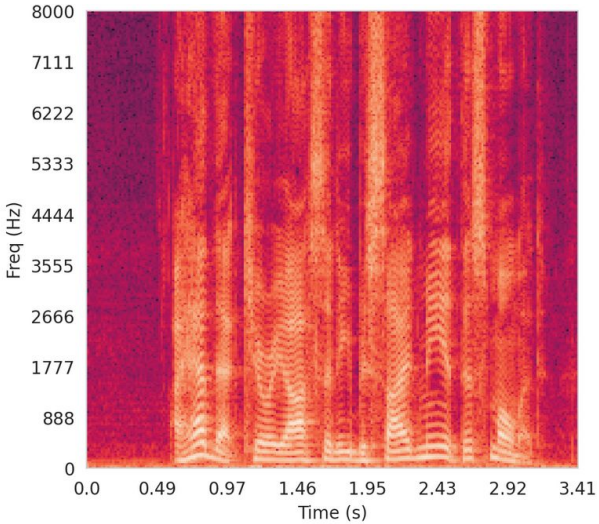

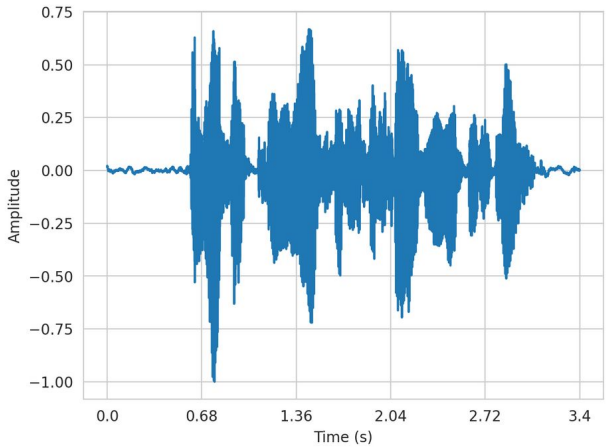

Mel Scale

Humans perceive sound on a log-scale. For human ear:

- 500 Hz << 600 Hz
- but 5000 Hz ~= 5100 Hz

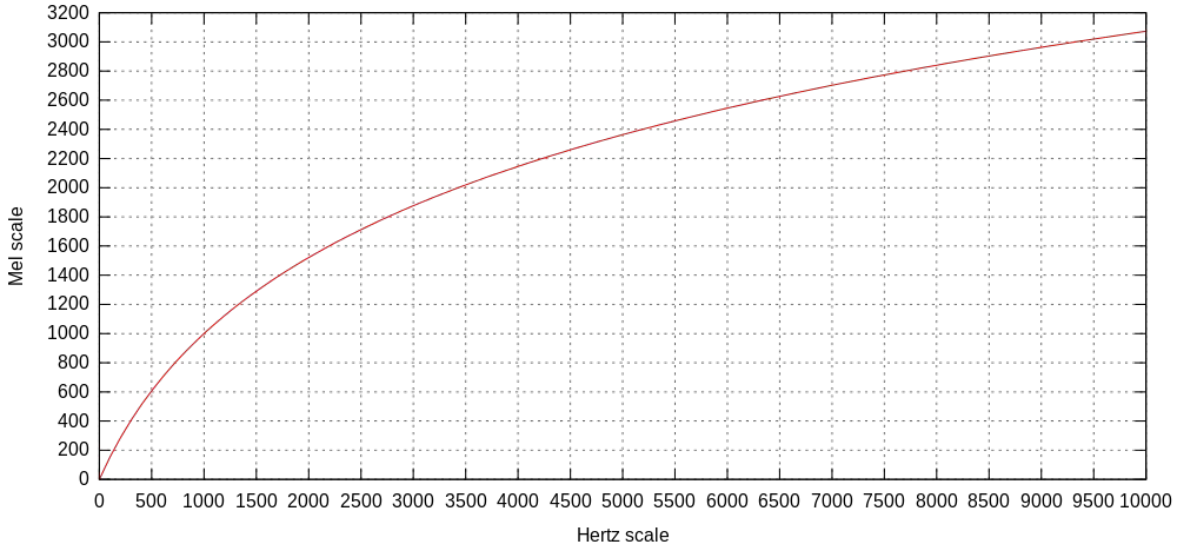

Mel Spectrogram

Many deep learning models use STFT in mel-scale instead. Note that this is a **lossy transform**

Weighted split on buckets

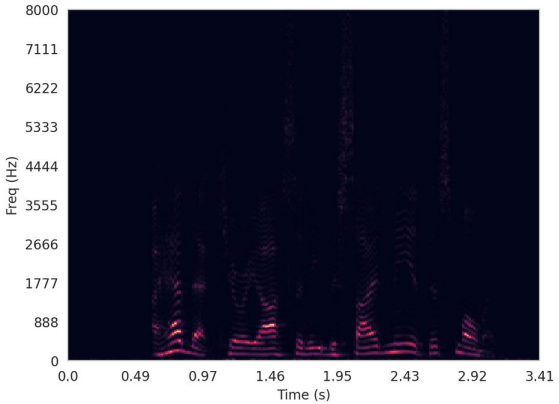

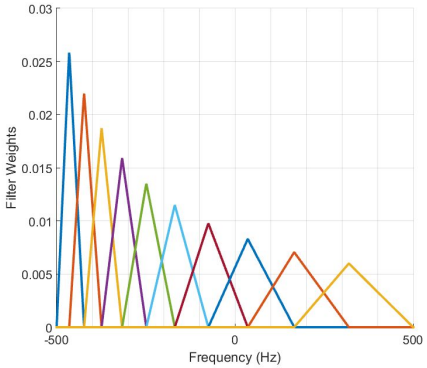

After log

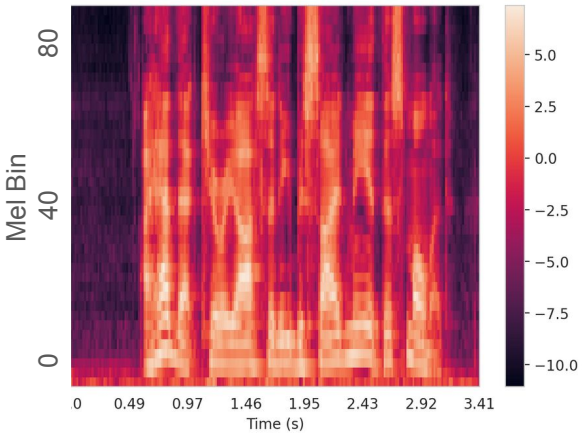

Implementation (librosa): https://librosa.org/doc/latest/generated/librosa.feature.melspectrogram.html#librosa.feature.melspectrogram

Mel Spectrogram

Remark: MelSpec is about making the **Y-axis look in log-scale**. Not about taking the log of the elements in the spectrogram!

Weighted split on buckets

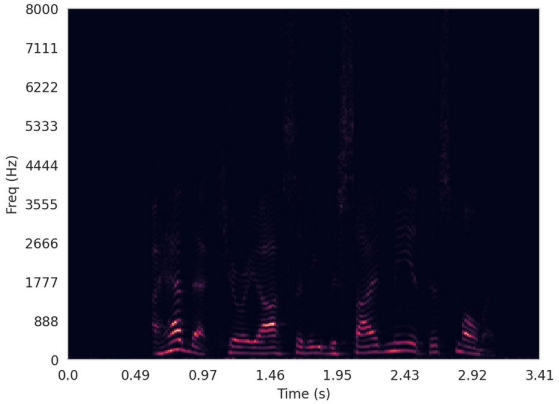

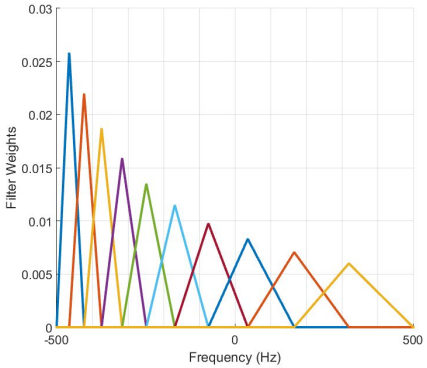

After log

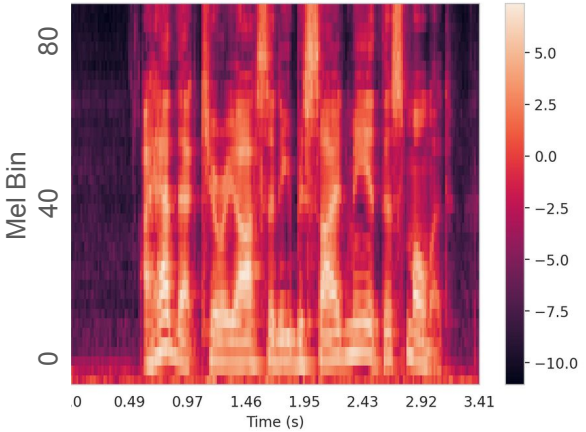

Tasks Overview and Voice Assistant

KWS

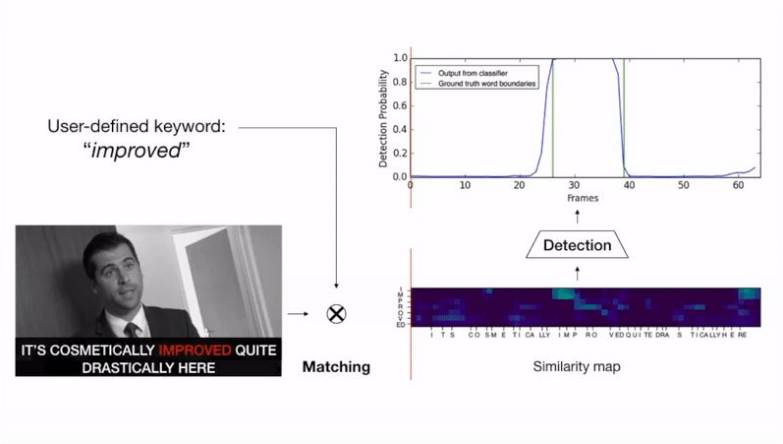

Momeni, Liliane, et al. "Seeing wake words: Audio-visual keyword spotting." arXiv preprint arXiv:2009.01225 (2020).

KWS is just a classification task.

Given a short signal with only one word, model is trained to distinguish keyword (label 1) from any other word (label 0)

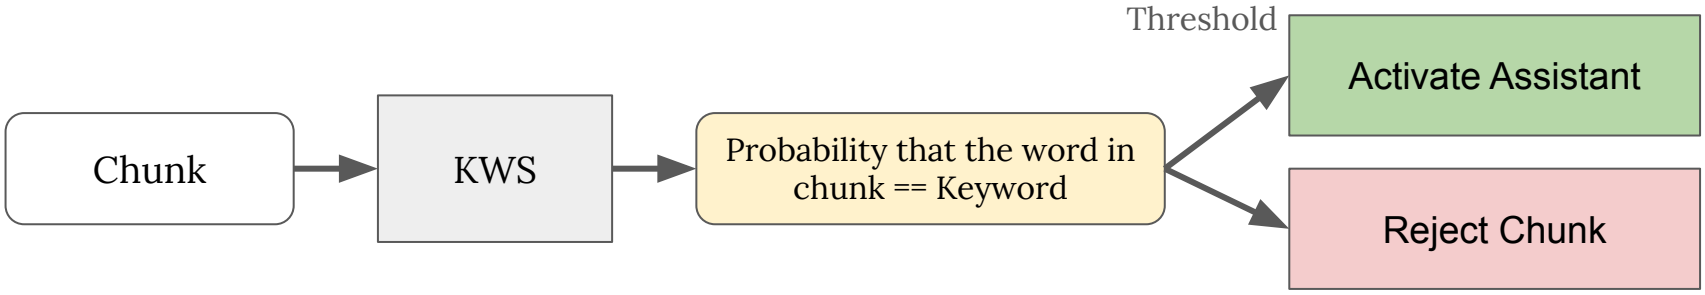

KWS

Let’s see how to train such a model in PyTorch

Automatic Speech Recognition (Speech To Text)

Generally, there 3 designs for ASR models

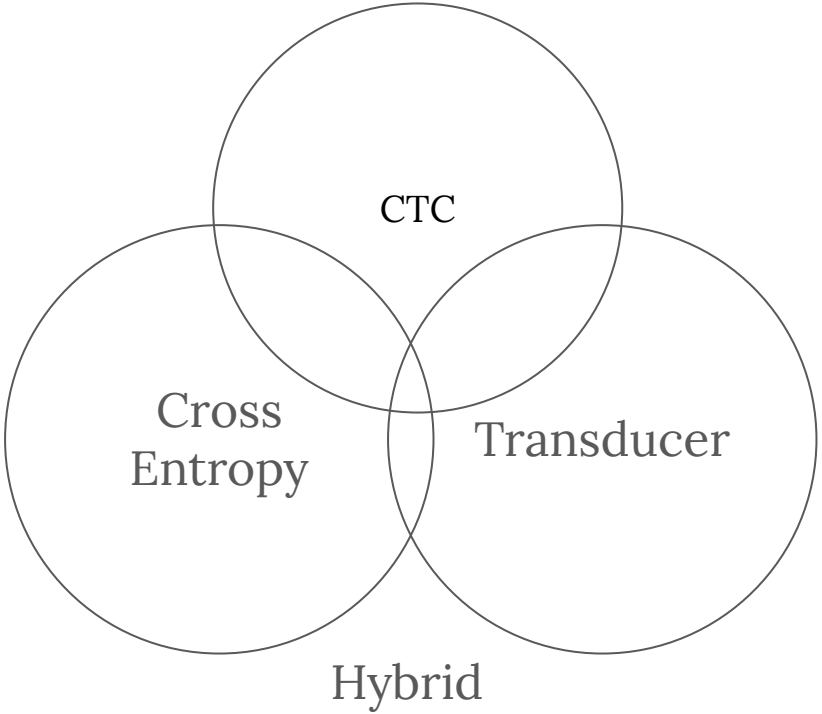

Automatic Speech Recognition (Speech To Text)

Why should you care? Each type has its own inference algorithms. Each algorithm also has its own sub-versions which may affect quality severely.

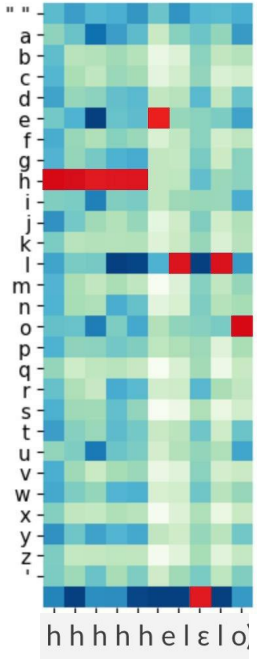

For example, in CTC, your model returns such probability matrix *time_step x vocabulary_size*

You can decode it into text via several ways:

- Greedy Search: Fast, but the quality suffer
- Beam Search: Slow, better quality
- Beam Search with Language Model: Even slower, remarkably better quality

Connectionist Temporal Classification (CTC)

In CTC, you have a vocabulary of size V: all possible tokens (in a simple case, letters from the alphabet and space: “ “)

Given speech input, the model returns the tensor TxV. Where each row i represents the probability that corresponding token is said at timestep i.

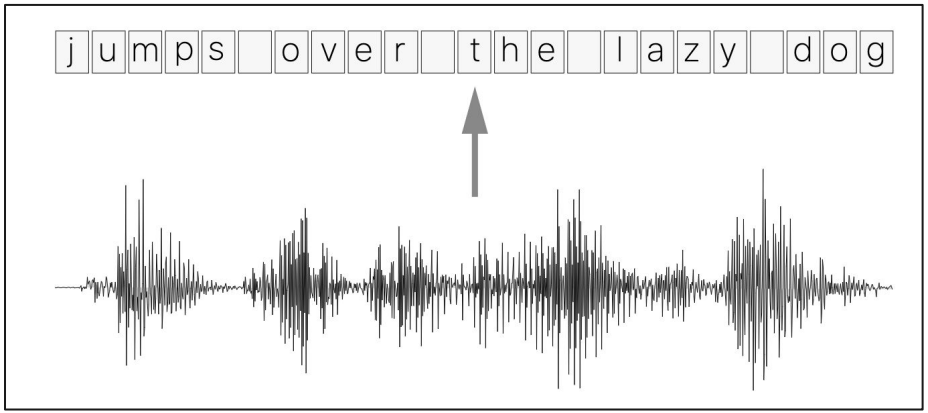

Typical encoder

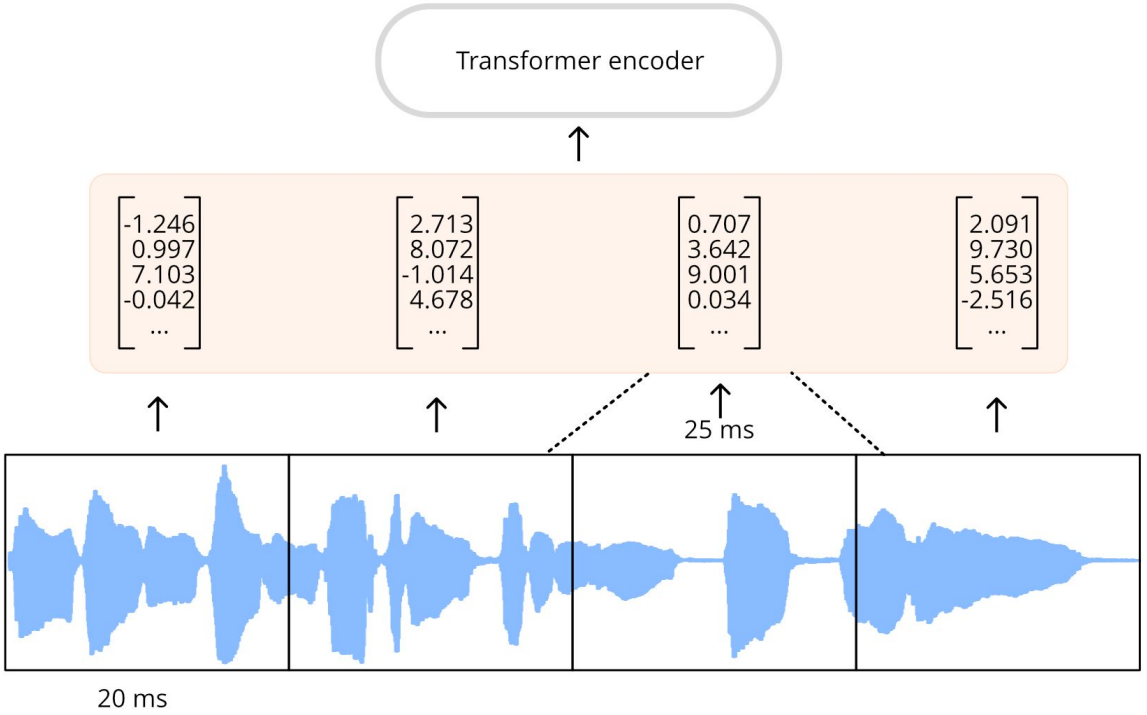

- Input audio is chunked (e.g. 20 ms)
- CTC outputs **one phoneme or character per chunk**

Image source: https://huggingface.co/learn/audio-course/chapter3/ctc

Connectionist Temporal Classification (CTC)

- Split input into frames
- Classify each frame into num_letters (|V|) classes
- Merge consecutive letters

Problem:
cannot distinguish some words:
He**ll**o vs Helo

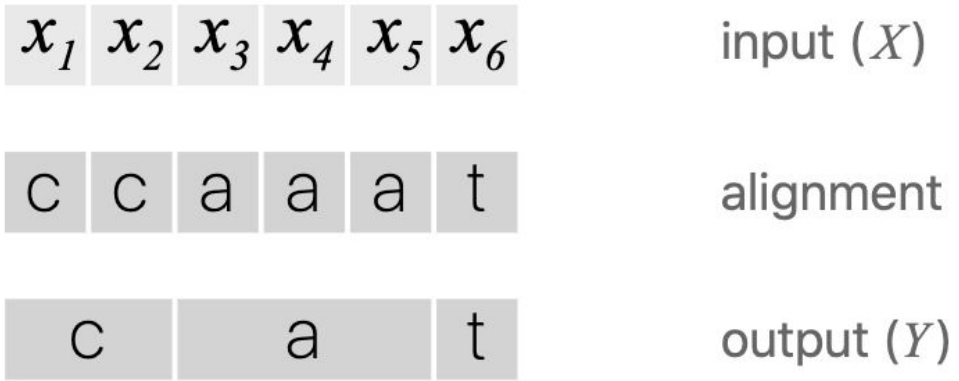

Connectionist Temporal Classification (CTC)

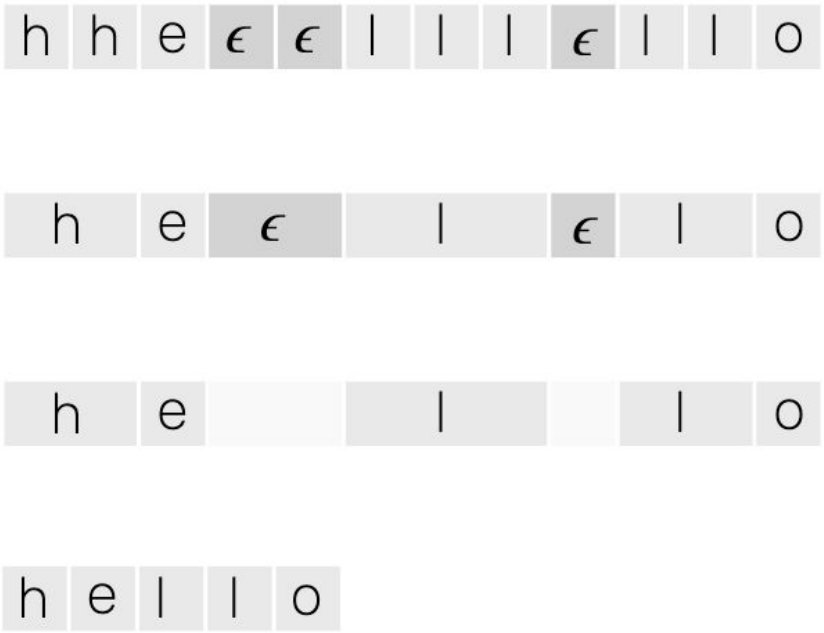

Solution: Add special token

First, merge repeat characters.

Then, remove any ϵ tokens.

The remaining characters are the output.

Connectionist Temporal Classification (CTC)

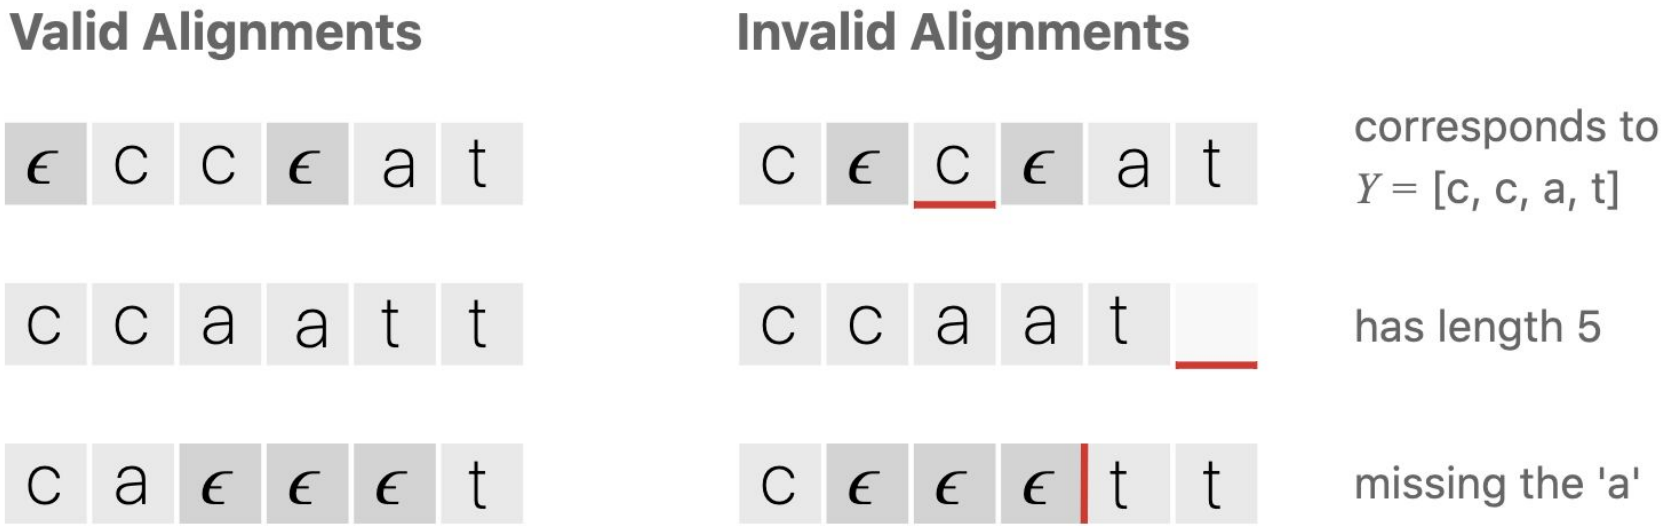

CTC Loss Function

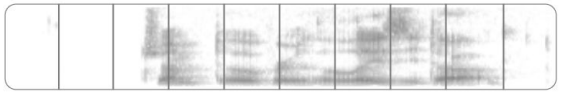

We start with an input sequence, like a spectrogram of audio.

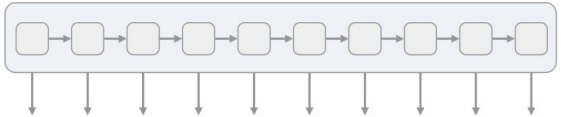

The input is fed into an RNN, for example.

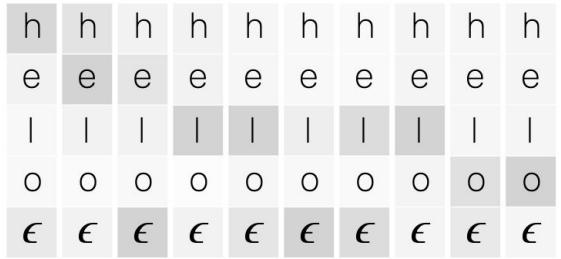

The network gives $p_t(a \mid X)$, a distribution over the outputs {h, e, l, o, ϵ} for each input step.

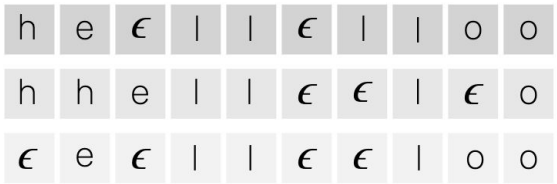

With the per time-step output distribution, we compute the probability of different sequences

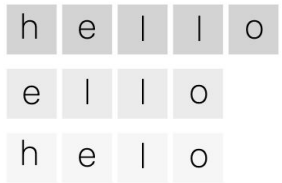

By marginalizing over alignments, we get a distribution over outputs.

CTC Loss Function

$$p(Y \mid X) = \sum_{A \in \mathcal{A}_{X,Y}} \prod_{t=1}^{T} p_t(a_t \mid X)$$

The CTC conditional **probability** | **marginalizes** over the set of valid alignments | computing the **probability** for a single alignment step-by-step.

Can be calculated effectively
via **dynamic programming** algorithm

$$p(\text{prefix} = \text{'cat'})_T = p(\text{prefix} = \text{'cat'}, \text{last\_char} = \text{'t'})_T + p(\text{prefix} = \text{'cat'}, \text{last\_char} = \epsilon)_T$$

$$p(\text{prefix} = \text{'cat'}, \text{last\_char} = \text{'t'})_T = p(\text{prefix} = \text{'ca'}, \text{last\_char} = \text{'a'})_{T-1} \cdot p(\text{'t'})_T + p(\text{prefix} = \text{'ca'}, \text{last\_char} = \epsilon)_{T-1} \cdot p(\text{'t'})_T + p(\text{prefix} = \text{'ca '}, \text{last\_char} = \text{'t'})_{T-1} \cdot p(\text{'t'})_T$$

CTC Loss Function

In PyTorch, everything is simple:

```
from torch.nn import CTCLoss

criterion = CTCLoss()
loss = criterion(
    log_probs=log_probs, # T x B x |V|
    targets=text_encoded, # B x S (length of tokenized text)
    input_lengths=log_probs_length,# B x T (length without padding)
    target_lengths=text_encoded_length, # B x S (length without padding)
)
```

Deep Speech (2014)

Input: Spectrogram

Model: MLP + Bi-LSTM

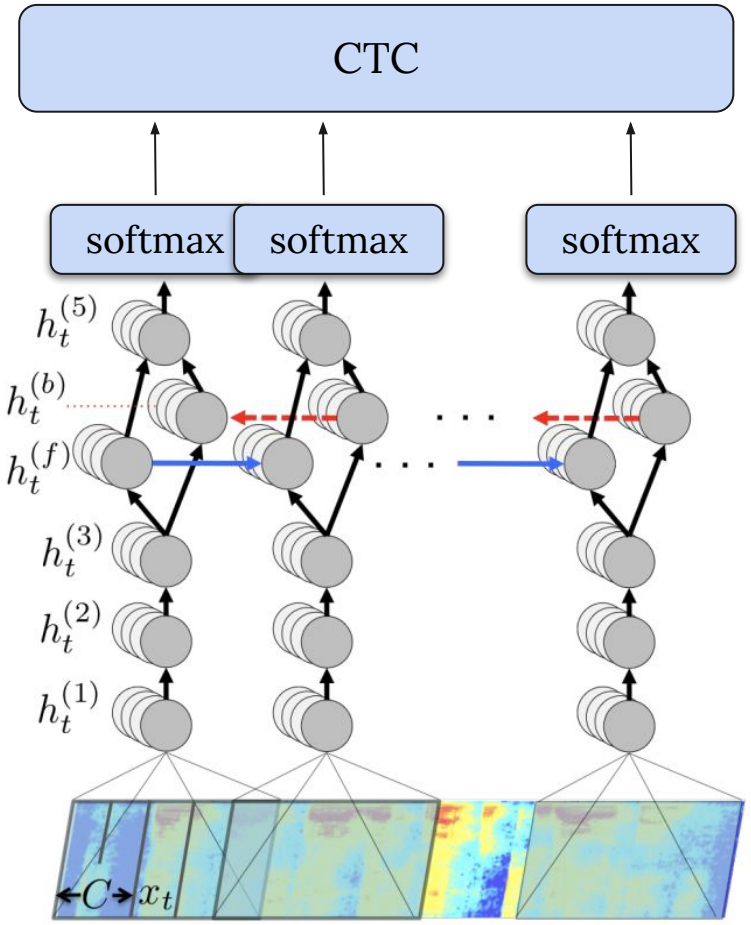

Deep Speech: Scaling up end-to-end speech recognition (2014)

Deep Speech 2 (2015)

Input: Spectrogram

Model: CNN + Bi-LSTM

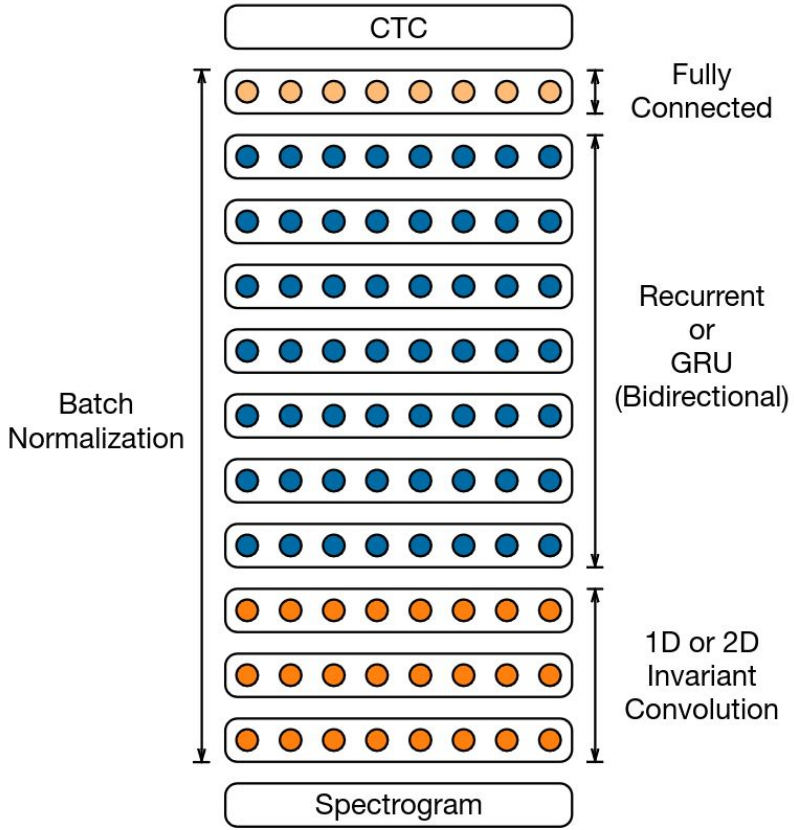

Deep Speech 2: End-to-End Speech Recognition in English and Mandarin (2015)

Conformer (2020)

Input: Spectrogram

Model: Modified Transformer

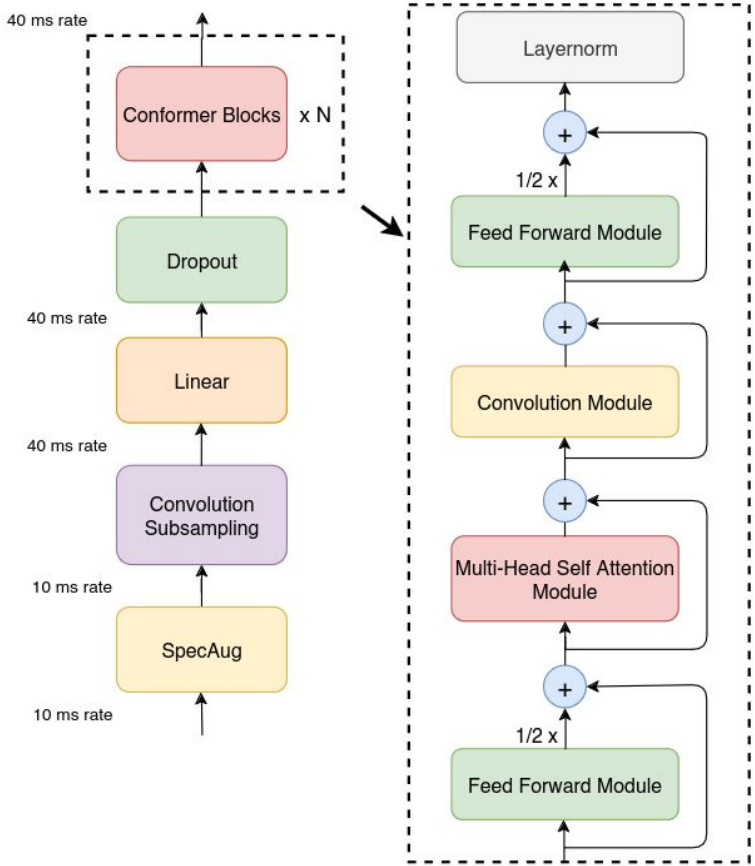

[Conformer: Convolution-augmented Transformer for Speech Recognition (Interspeech 2020)]

Conformer (2020) - Convolution Module

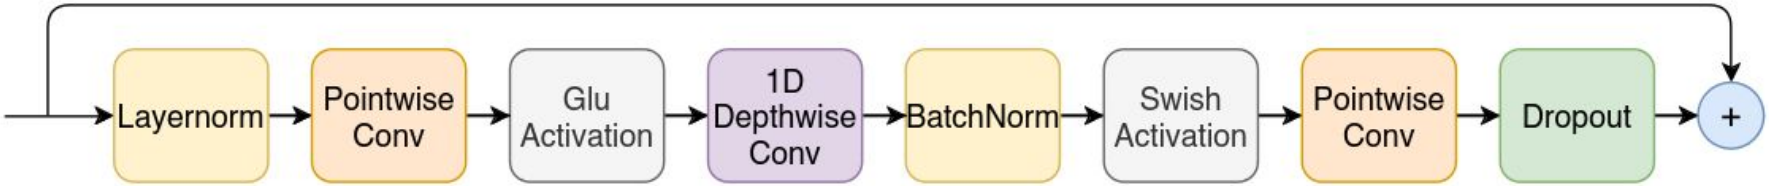

Depthwise Convolution

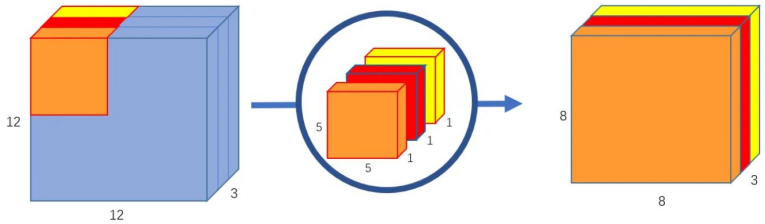

Pointwise Convolution

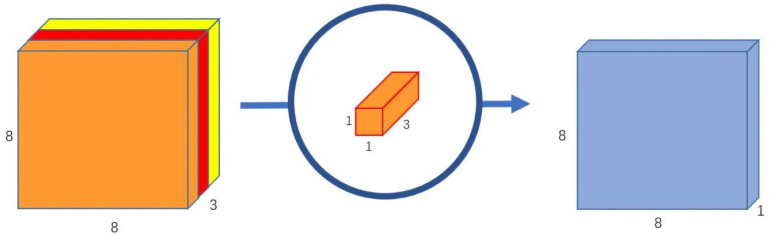

Far-field ASR

Reverberation

Interferences

Self-sound

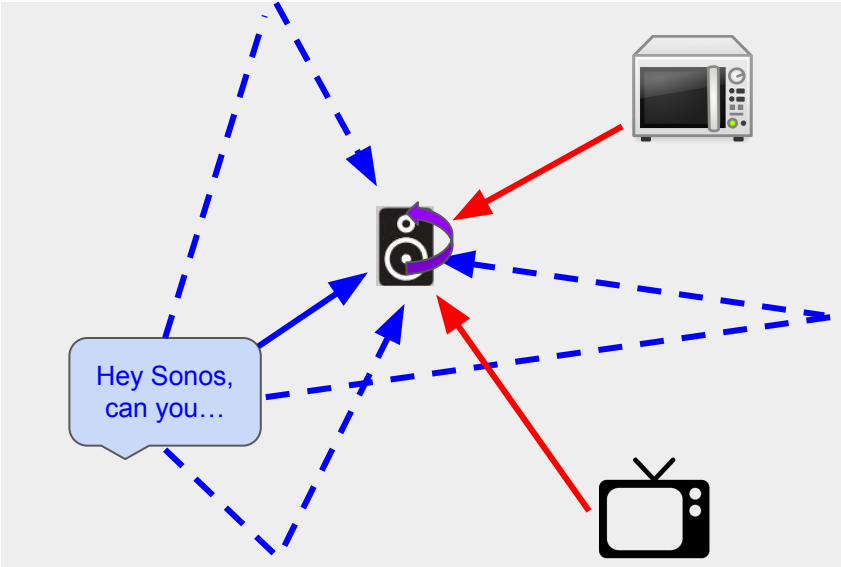

Dereverberation

Spectral masking, beamforming

Adaptive filtering

Propagation to microphone(s) can be modeled with room impulse response(s)

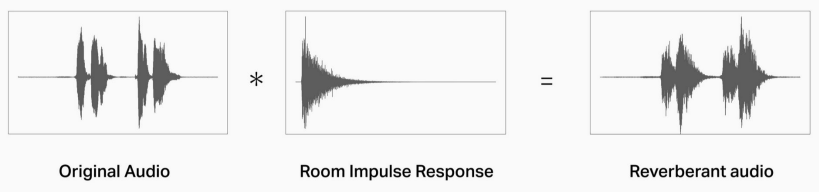

Robust performance with data augmentation

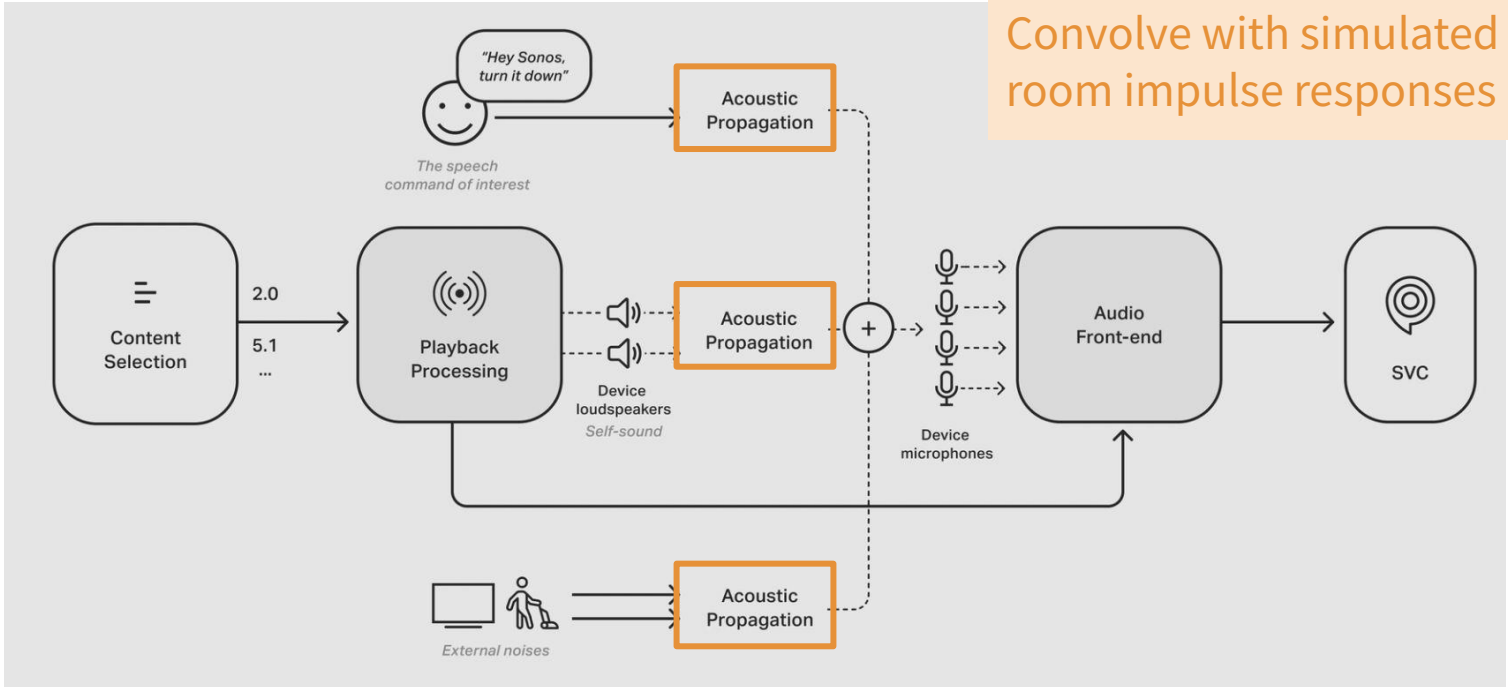

Convolve with simulated room impulse responses (RIRs).

[Data Augmentation for SVC]

Simulating room impulse responses: https://github.com/LCAV/pyroomacoustics

More data augmentation

- Changing speed
- Changing pitch
- SpecAugment (time and frequency mapping)

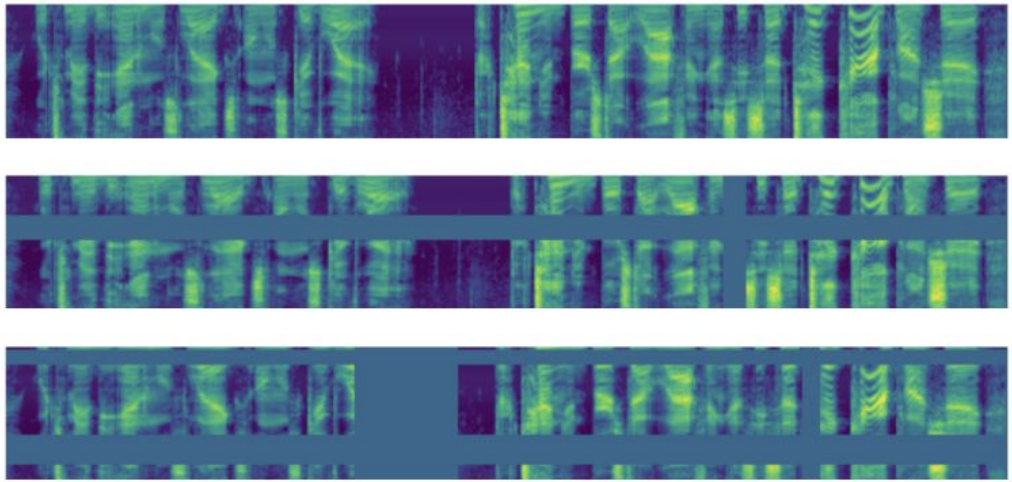

Example code: https://speechbrain.readthedocs.io/en/latest/tutorials/preprocessing/speech-augmentation.html

Automatic Speech Recognition (Speech To Text)

You can check HuggingFace Open ASR LeaderBoard and choose the model that fits your requirements

Leaderboard | Metrics | Request a model here!

model	Average WER	RTFx	AMI	Earnings22	Gigaspeech	LS Clean	LS Other
nvidia/canary-1b	6.5	235.34	13.9	12.19	10.12	1.48	2.93
nvidia/parakeet-tdt-1.1b	7.01	2390.61	15.87	14.49	9.52	1.4	2.6
nvidia/parakeet-rnnt-1.1b	7.12	2053.15	17.01	13.94	9.89	1.45	2.5
nvidia/parakeet-ctc-1.1b	7.4	2728.52	15.67	13.75	10.28	1.83	3.51
openai/whisper-large-v3	7.44	145.51	15.95	11.29	10.02	2.01	3.91
nvidia/parakeet-tdt_ctc-110m	7.49	5345.14	15.89	12.37	10.52	2.4	5.22
nvidia/parakeet-rnnt-0.6b	7.5	2815.72	17.4	14.66	10.01	1.62	3.02
distil-whisper/distil-large-v3	7.52	214.42	15.16	11.79	10.08	2.54	5.19
nvidia/parakeet-ctc-0.6b	7.69	4281.53	16.46	14.26	10.39	1.88	3.8

Text To Speech (TTS)

For TTS, most systems are divided into 2 parts: (i) Text-To-Intermediate Representation and (ii) Vocoder (Representation-To-Waveform)

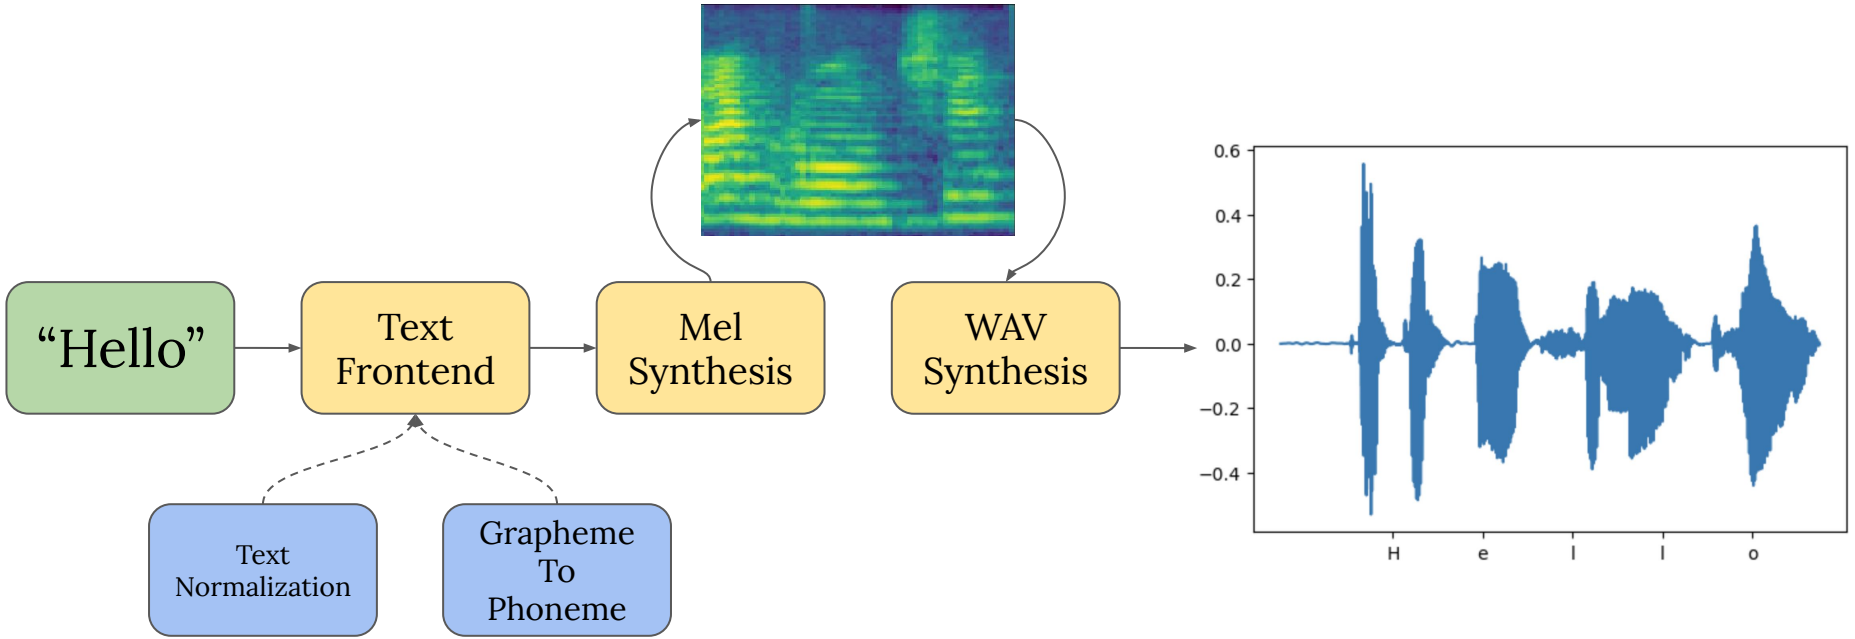

Text To Speech - Wav Synthesis

Due to compression, Mel Spectrogram is not invertible (**lossy** transform)

We can use Generative Adversarial Networks to generate raw audio from Mel Spectrogram.

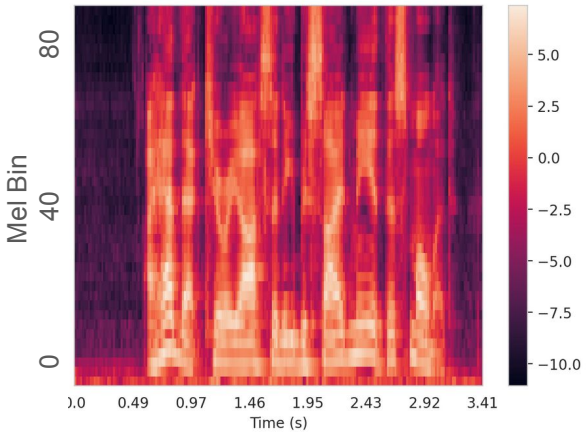

GAN

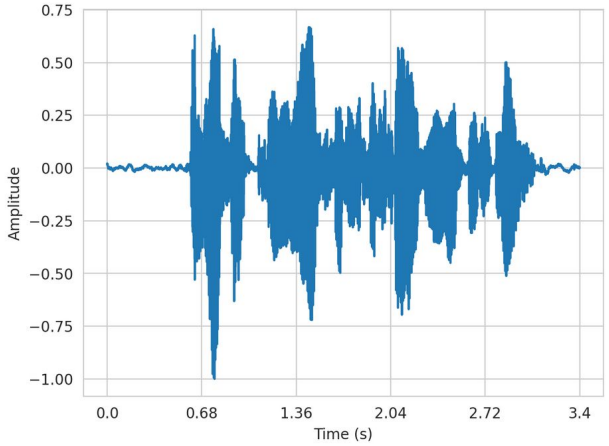

Remark: GANs

GAN consists of 2 networks: **Generator** and **Discriminator**

Discriminator solves binary classification task.

Given real and fake examples it predicts are they real or fake

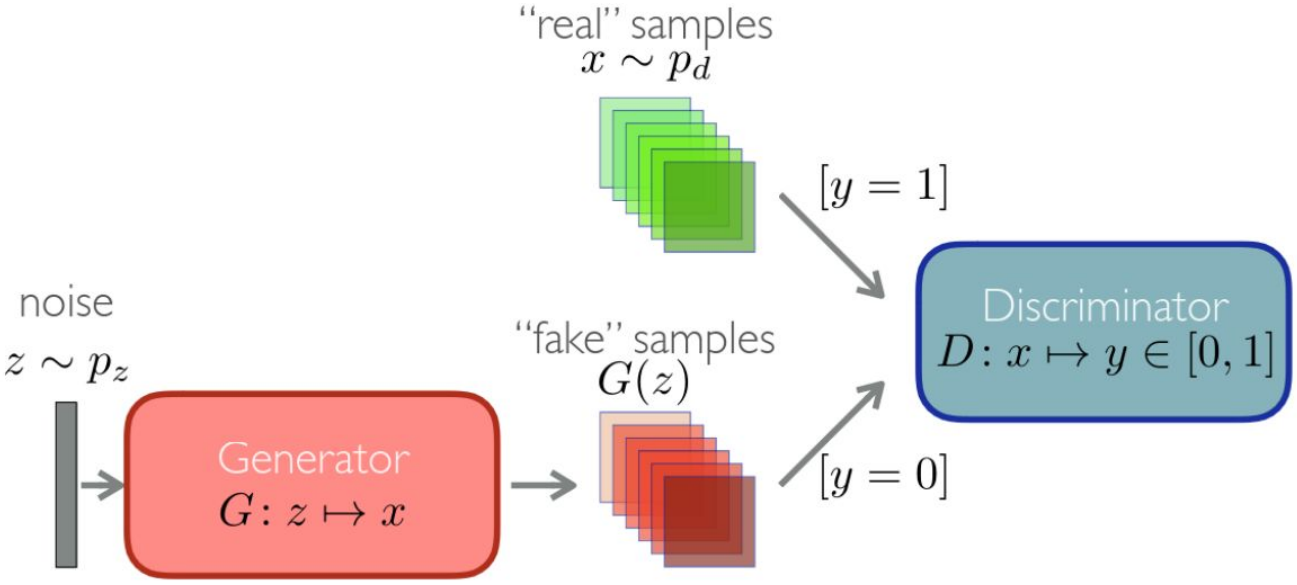

Remark: GANs

GAN consists of 2 networks: **Generator** and **Discriminator**

Generator tries to generate realistic examples and fool the discriminator

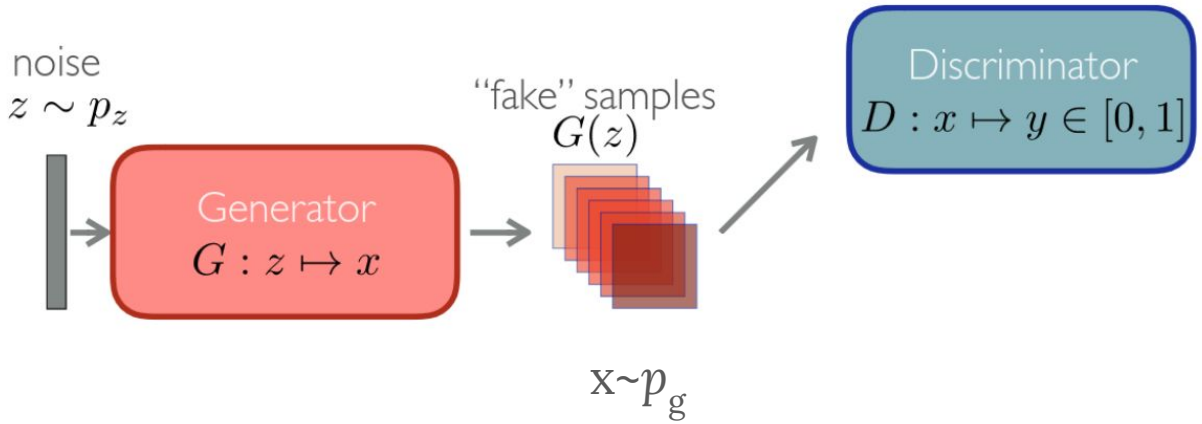

Remark: GANs

There are alternative Loss Functions:
See LSGAN and WGAN

Two Networks play 2 players game:

$$\min_G \max_D \mathop{\mathbb{E}}_{x \sim p_d} [\log D(x)] + \mathop{\mathbb{E}}_{z \sim p_z} [\log(1 - D(G(z)))]$$

$$\mathscr{L}_D(G, D) = \max_D \mathop{\mathbb{E}}_{x \sim p_d} [\log D(x)] + \mathop{\mathbb{E}}_{z \sim p_z} [\log(1 - D(G(z)))]$$

$$\mathscr{L}_G(G, D) = \min_G \mathop{\mathbb{E}}_{z \sim p_z} [\log(1 - D(G(z)))]$$

$$:= (\text{in practice}) \max_G \mathop{\mathbb{E}}_{z \sim p_z} [\log(D(G(z)))]$$

Remark: Training GANs

Train via alternating algorithm: Train Discriminator, then Generator

```
g_outputs = model.generate(**batch)
batch.update(g_outputs)

D_optimizer.zero_grad()

d_outputs = model.discriminate(generated_x=batch["generated_x"].detach(),
                               real_x=batch["real_x"])
batch.update(d_outputs)

D_loss = discriminator_criterion(**batch)
D_loss.backward()
D_optimizer.step()

G_optimizer.zero_grad()
d_outputs = model.discriminate(**batch)
batch.update(d_outputs)
G_loss = generator_criterion(**batch)
G_loss.backward()
```

Remark: Conditional GANs (CGAN)

Conditional GANs have some other input to condition on apart from noise z

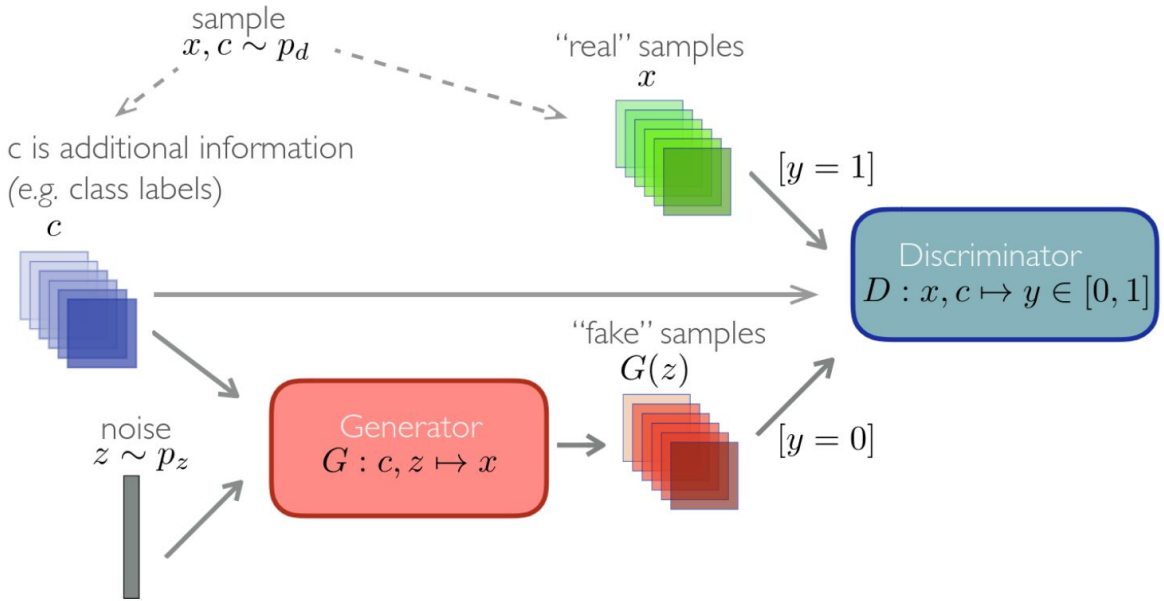

Remark: Deep Convolutional GAN (DCGAN)

Classic architecture DCGAN (Deep-Convolutional GAN)

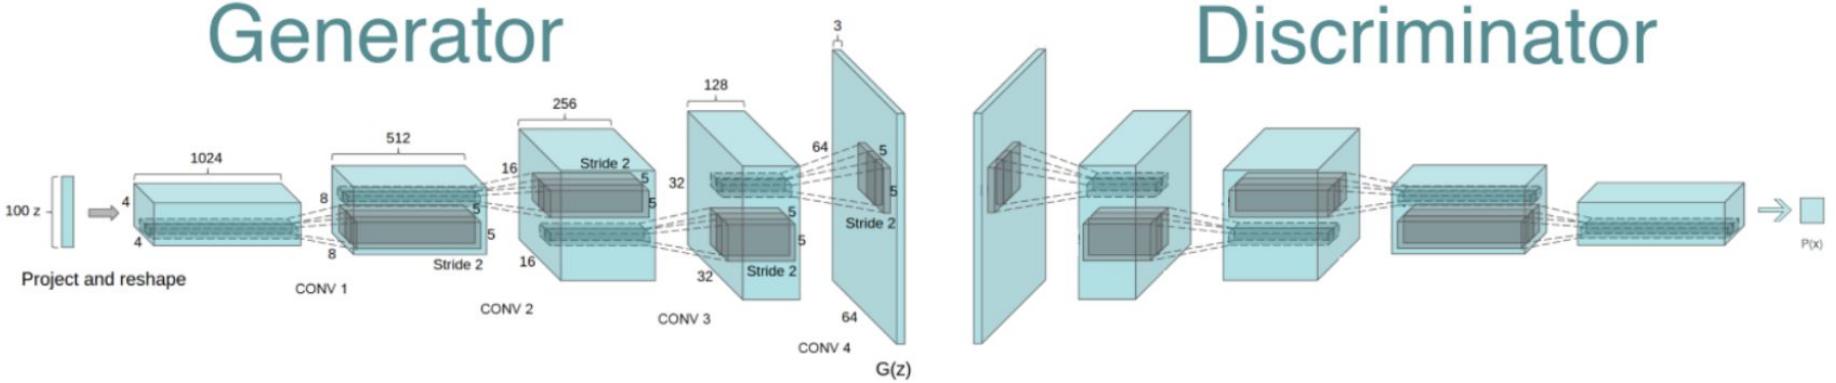

HiFi-GAN (2020)

Generator converts Mel Spectrogram to a Waveform.

Many ResBlock with different Kernels/Dilations to model various receptive fields

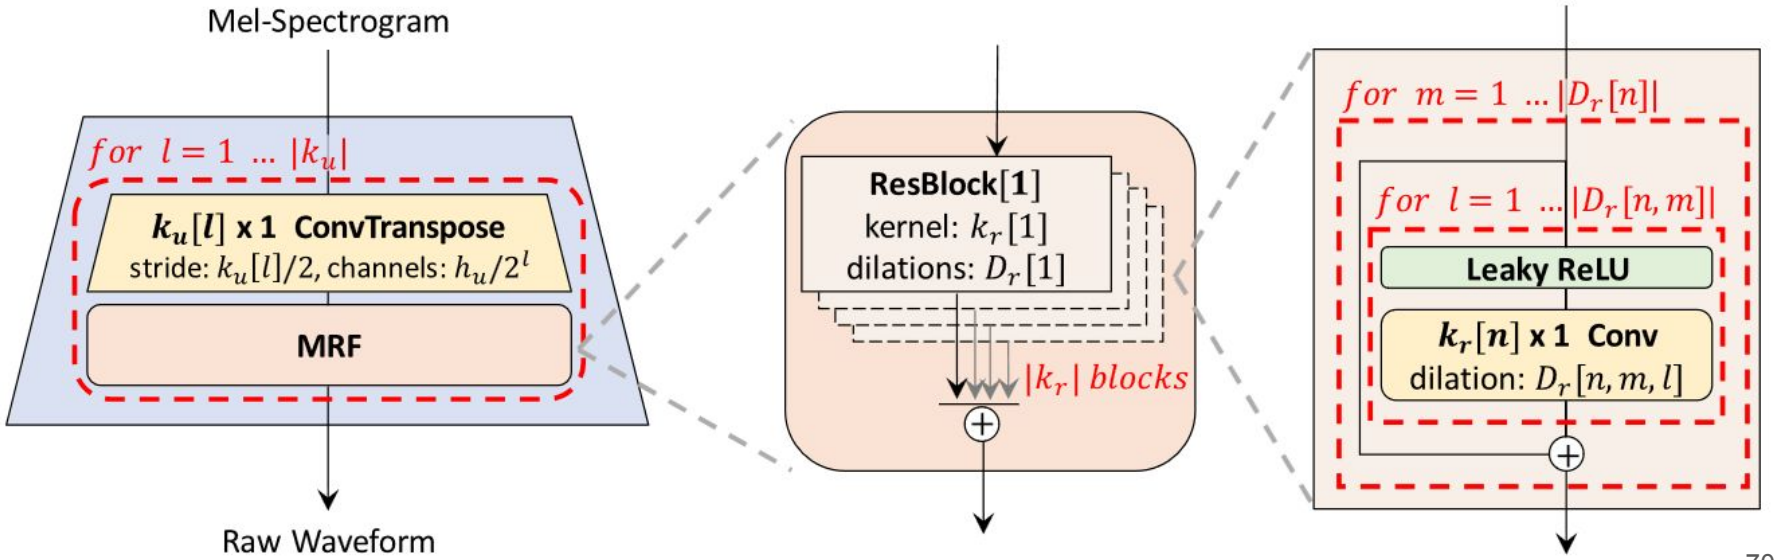

HiFi-GAN: Generative Adversarial Networks for Efficient and High Fidelity Speech Synthesis (NIPS 2020)

HiFi-GAN (2020)

Multi-Period Discriminator analysis audio on different periods p [2, 3, 5, 7, 11]

This allows to look at diverse periodic patterns in the audio (recall that audio is a sum of sinusoids)

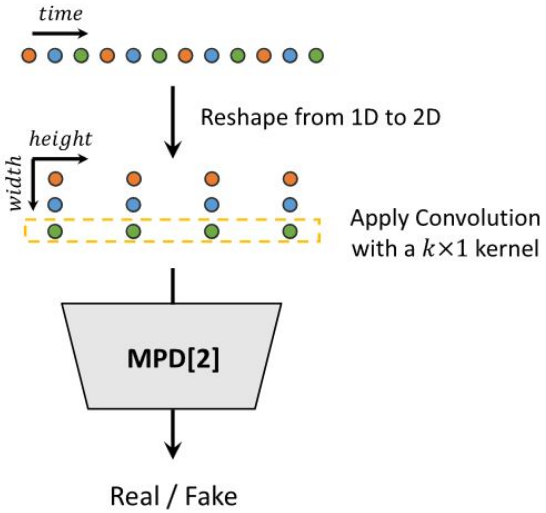

HiFi-GAN: Generative Adversarial Networks for Efficient and High Fidelity Speech Synthesis (NIPS 2020)

HiFi-GAN (2020)

Multi-Scale Discriminator analysis audio on different levels with different receptive fields

This allows to capture consecutive patterns and long-term dependencies

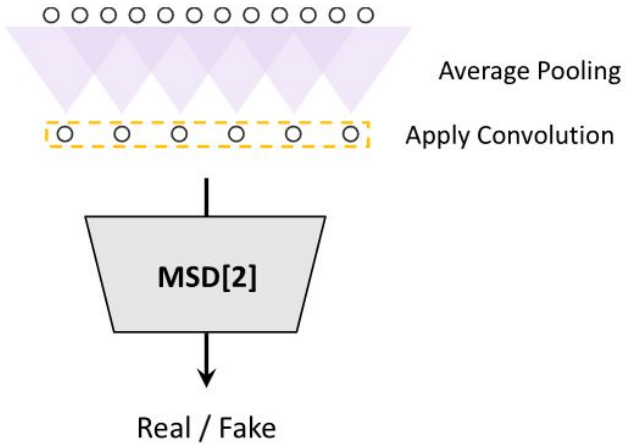

HiFi-GAN: Generative Adversarial Networks for Efficient and High Fidelity Speech Synthesis (NIPS 2020)

Voice Conversion (VC)

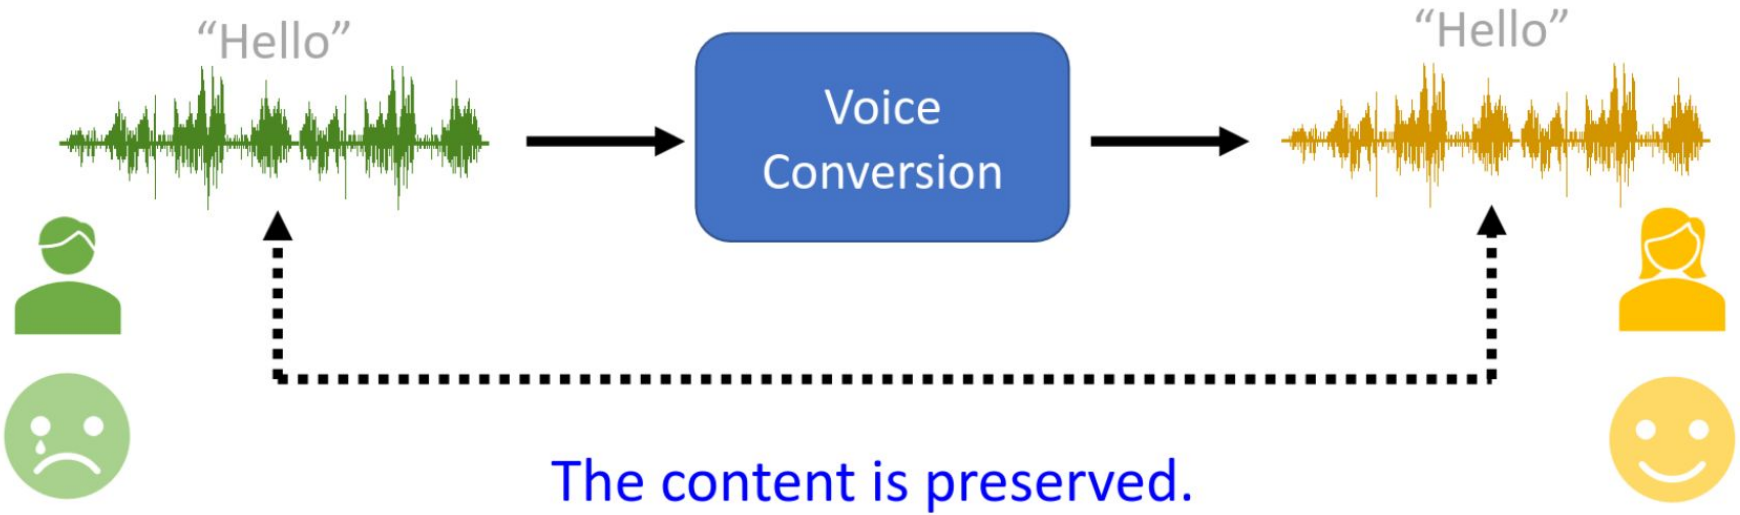

Voice Cloning – similar stuff, but the speaker is preserved, content is changed

Voice Conversion (VC)

A common design approach is to have 2 networks. One to extract content information (what speaker says) and another one to extract speaker information (tembre, pitch, etc.). The outputs are combined to get a new audio with the content of the first speaker and identity of the second one

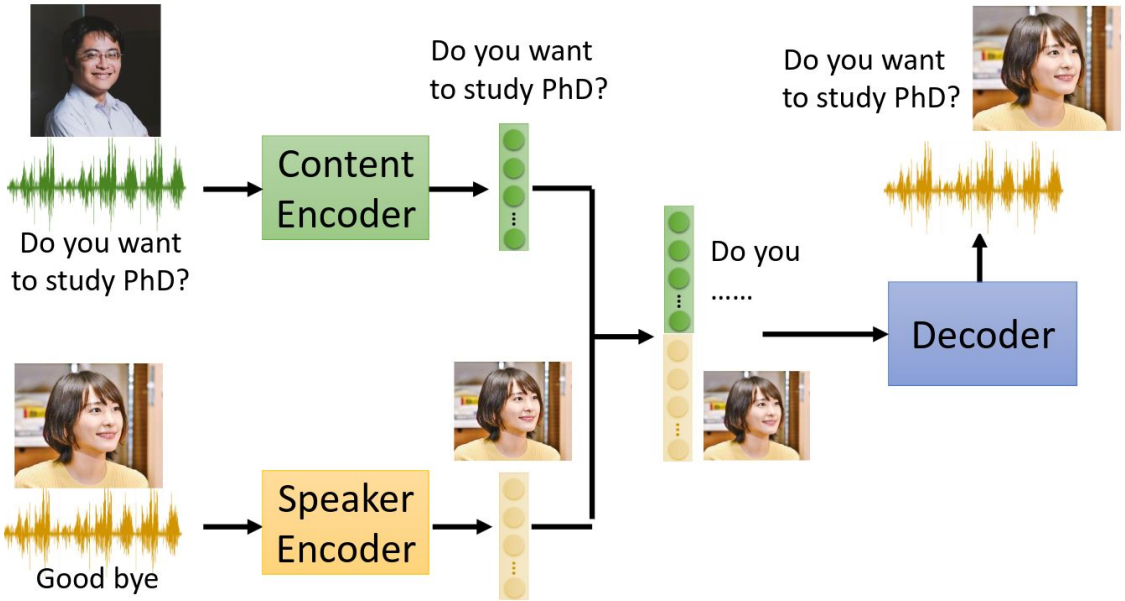

HiFi-VC (2023)

Combines HiFi-GAN from TTS, Encoder from ASR, VAE from Voice Biometry

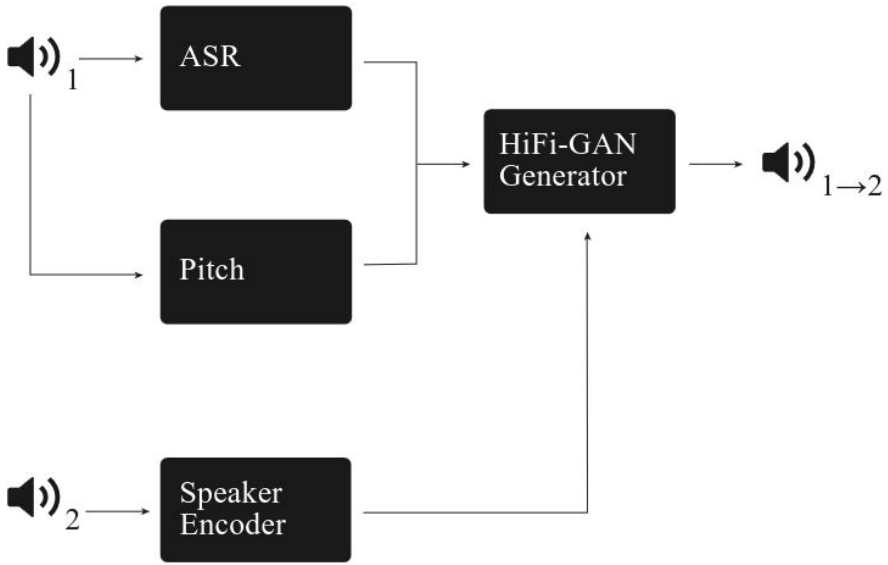

HiFi-VC: High Quality ASR-Based Voice Conversion (Interspeech 2023)

HiFi-VC (2023)

Combines HiFi-GAN from TTS, Conformer from ASR, VAE from Voice Biometry

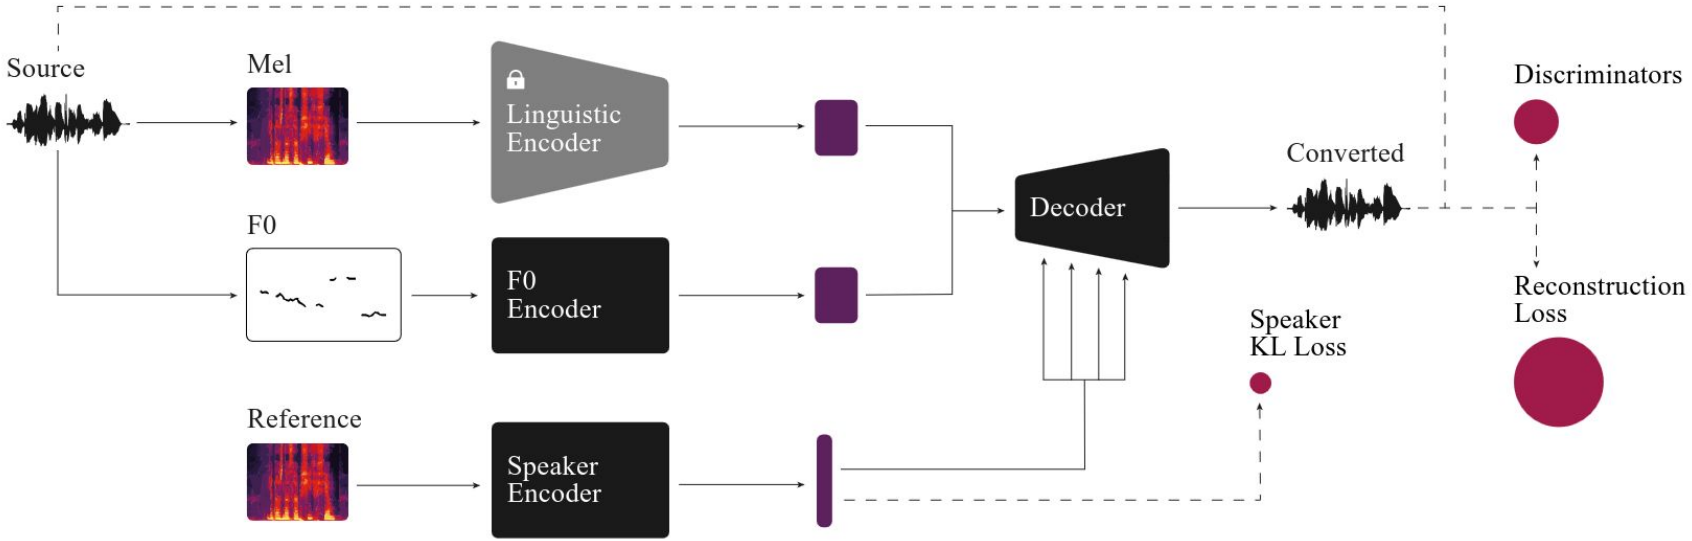

Want to know more about DLA?

See our HSE DLA Course and LauzHack Deep Learning Bootcamp

https://github.com/markovka17/dla

https://github.com/LauzHack/deep-learning-bootcamp

Let’s Start Building An Assistant

The Voice Assistant Code

The code for the assistant is available via the link below. We will go through it right now.

The Voice Assistant Pipeline: KWS and Streaming

We start with a KWS system that will activate the assistant (like “Hey, Siri” or “Ok, Google”).

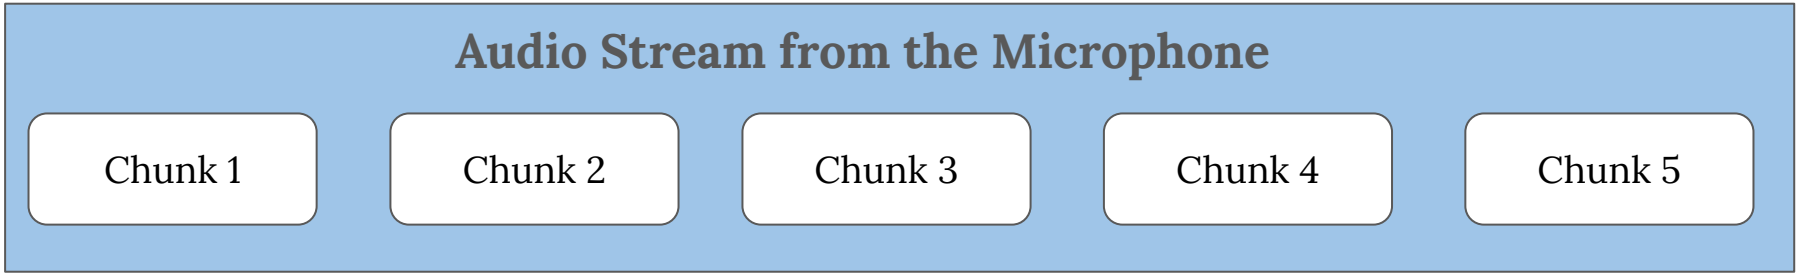

We will process chunks successively:

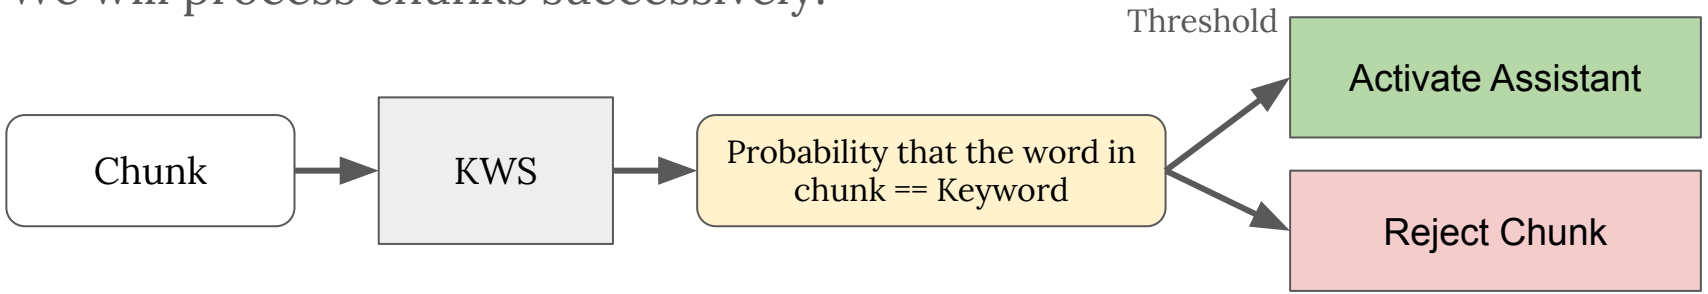

Reading from the microphone via PyTorch

To ensure fast processing, it is better to separate the main process where we apply our models from the process that reads the stream. For this, we will use *multiprocessing lib* and two processes.

```
94  def init_stream(all_models, history, config):
95      ctx = mp.get_context("spawn")
96      manager = ctx.Manager()
97      chunk_queue = manager.Queue(maxsize=0)
98      stream_init_queue = manager.Queue(maxsize=0)
99      streaming_process = ctx.Process(
100         target=audio_stream,
101         args=(
102             stream_init_queue,
103             chunk_queue,
104             config.stream_reader.source,
105             config.stream_reader.format,
106             config.stream_reader.chunk_size,
107             config.stream_reader.sample_rate,
108             config.stream_reader.vad_limit,
109         ),
110     )
```

Reading from the microphone via PyTorch

The **Queue** object controls the communication between the processes. We will use it to send audio chunks from the microphone to the main process. We will also send control commands from the main process to the microphone one (“start recording”, “stop recording”)

```
49      streamer = StreamReader(src=source, format=format)
50      while init_queue.get() == SIG_START:  # main process sent the signal to start
51          streamer.add_basic_audio_stream(
52              frames_per_chunk=chunk_size, sample_rate=sample_rate
53          )
54          stream_iterator = streamer.stream(timeout=-1, backoff=1.0)
```

Reading from the microphone via PyTorch

Read from the microphone stream

```
48          print("Start audio streaming")
49          while True:
50              (chunk_,) = next(stream_iterator)
```

Sending chunk to the main process for further processing.

```
64
65              queue.put(chunk_data)
```

Sending command that we achieved recording stop criteria

```
89          queue.put(SIG_STOP)  # stop criteria for main process loop
90          print("End of the user input")
91          streamer.remove_stream(0)  # remove stream with id=0
```

Reading from the microphone via PyTorch

In the main process, we init the recording process, send the command to start the recording and get chunks from the recording process:

```
122     streaming_process.start()
123     stream_init_queue.put(SIG_START)  # sent signal to start
124     while True:
125         try:
126             while True:
127                 chunk = chunk_queue.get()
128                 if isinstance(chunk, int):  # stop criteria achieved
129                     break
130                 user_chunks.append(chunk)
```

Reading from the microphone via PyTorch

After processing the chunks (and using TTS to generate response), we send the signal to start recording again. In case we want to exit, we send exit code

```
156                # init next user input
157                user_chunks = []
158                stream_init_queue.put(SIG_START)
159            except KeyboardInterrupt:
160                break
161            except Exception as exc:
162                raise exc
163
164        # send signal to exit
165        stream_init_queue.put(SIG_STOP)
166        streaming_process.join()
167        stream_writer.close()
```

Join command ensures that we will wait for the end of the second process

The Voice Assistant Pipeline: KWS and Streaming

We start with a KWS system that will activate the assistant (like “Hey, Siri” or “Ok, Google).

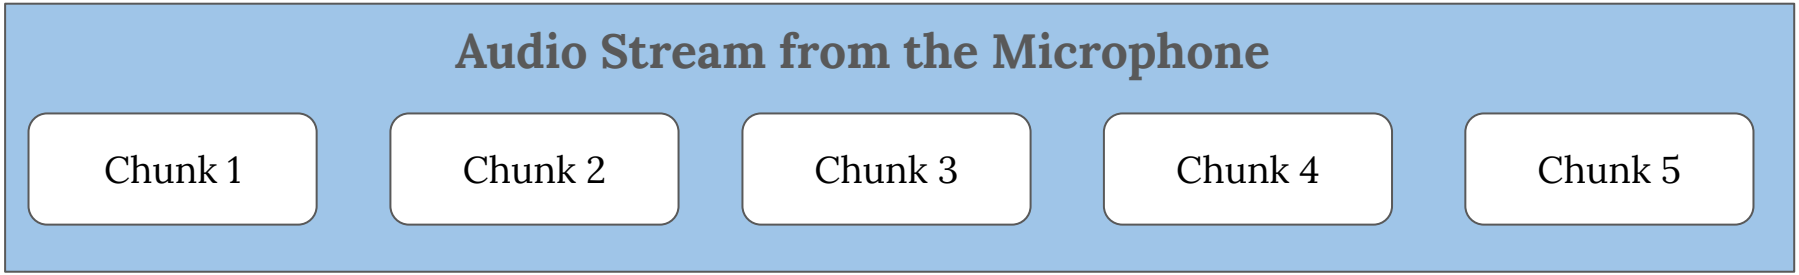

We will process chunks successively:

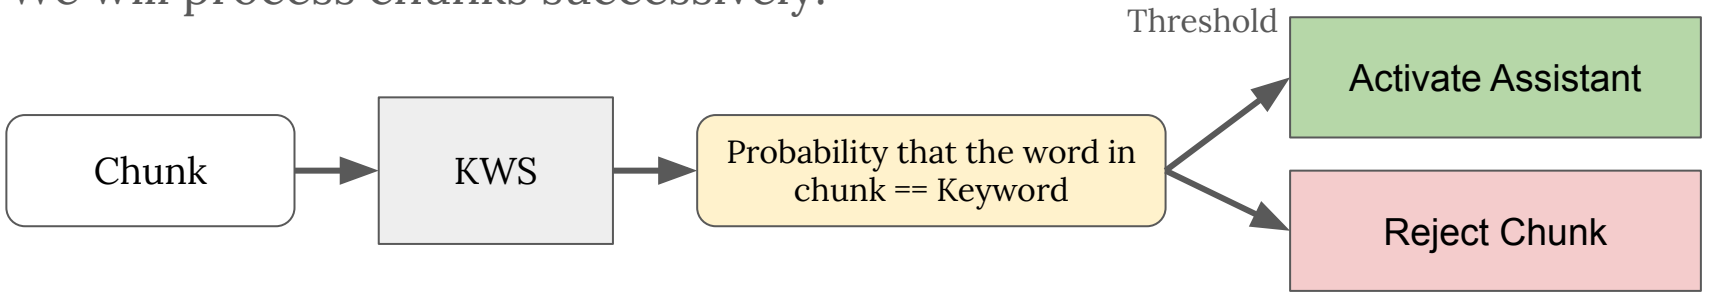

The Voice Assistant Pipeline: KWS and Streaming

Ideally, the streaming process should only record audio, but for the simplicity we add some checks like KWS in it too.

```
61          print("Start audio streaming")
62          while True:
63              (chunk_,) = next(stream_iterator)
64              chunk_data = chunk_.data
65              chunk_data = chunk_data.sum(dim=-1)  # to mono
66              chunk_data = chunk_data.view(1, -1)
67
68              if not query_started:
69                  # check that the query is started
70                  with torch.inference_mode():
71                      keyword_proba = kws(chunk_data)
72                  if keyword_proba > KWS_THRESHOLD:
73                      print(f"Keyword Detected: {keyword_proba}")
74                      query_started = True
75                  else:
76                      continue
77
78              queue.put(chunk_data)
```

The Voice Assistant Pipeline: VAD

After the Assistant is **activated**, we need to listen for the user's query. How to understand whether is it ended or not?

We will listen to all the chunks after assistant is activated

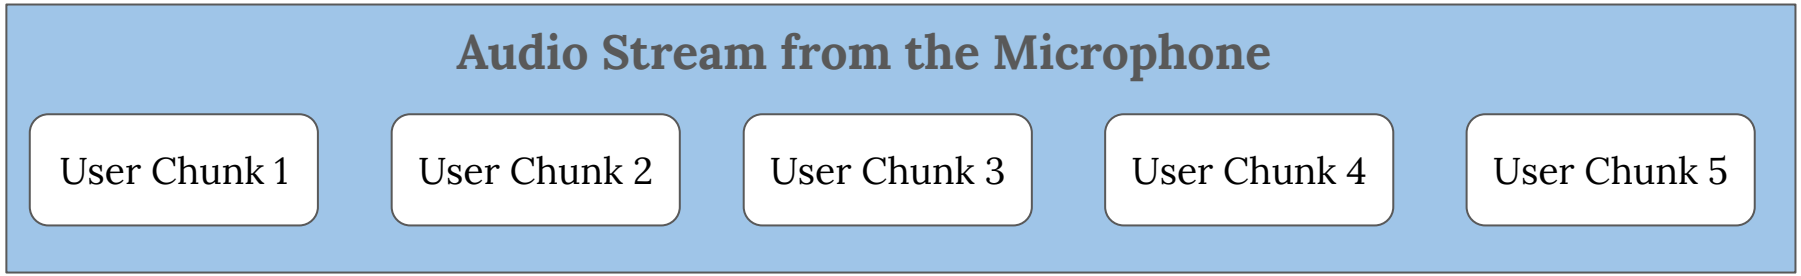

If the user is silent for several successive chunks, we will believe he ended the query. We can check silence via Voice Activity Detector

The Voice Assistant Pipeline: VAD

```
79                vad_check = vad_checking(chunk_data, sample_rate=sample_rate
80                if vad_check:
81                    vad_counts += 1
82                else:
83                    vad_counts = 0
84                    user_talked = True
85                if user_talked and vad_counts >= vad_limit:
86                    break
87
88            queue.put(SIG_STOP)  # stop criteria for main process loop
89            print("End of the user input")
90            streamer.remove_stream(0)  # remove stream with id=0
91
```

The Voice Assistant Pipeline: Getting Audio Response

With the help of VAD, we got the whole user's query. Now we need to generate a response:

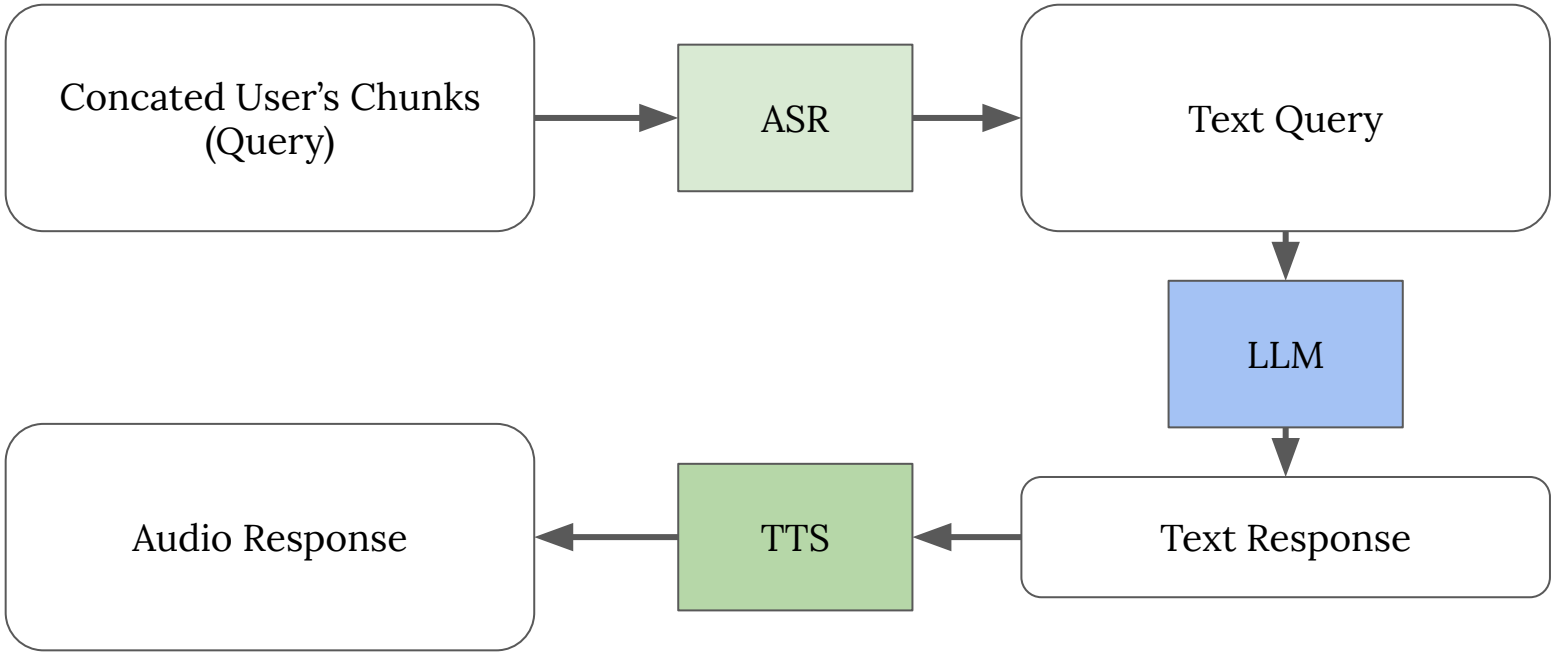

Remark: Can we avoid ASR and feed the audio directly?

Novel LLMs are capable of working with Audio directly. However, the pipeline is almost the same, the model just skips explicit text representation. Example from LLAMA-3

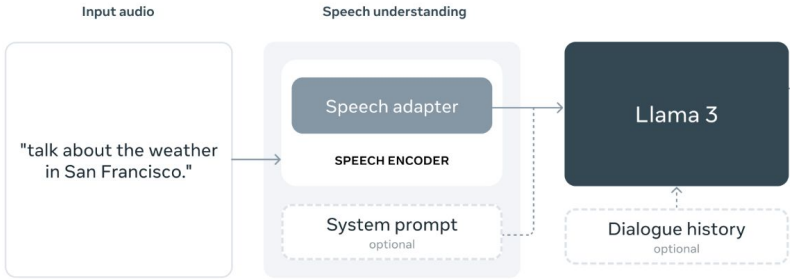

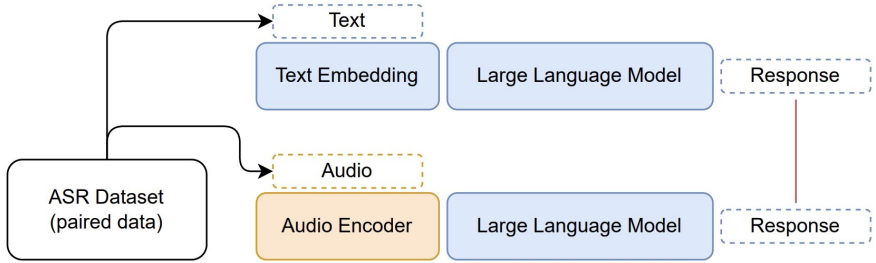

Figure 2: The aim is to create an end-to-end system that would generate the same response when being fed spoken (audio) input instead of its text version. The overall system should be invariant to the modality of the inputs containing the same semantic information.

AudioChatLlama: Towards General-Purpose Speech Abilities for LLMs (2024), The Llama 3 Herd of Models (2024)

Where to get models?

The first approach is to train a model yourself, like we did for KWS. The second option is to use pre-trained models

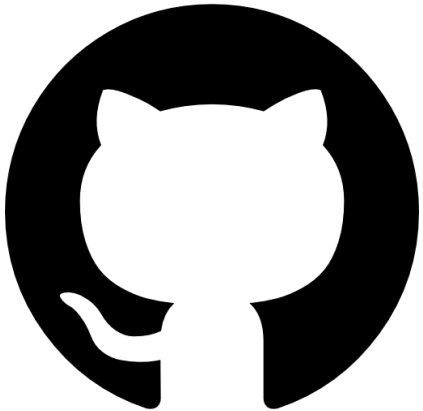

General Use: HuggingFace and GitHub contain many pre-trained models. For audio and many other domains.

Where to get models?

Audio-Specific: There are several libraries that provide pre-trained models specifically for audio domain (some of them have other domains too, but the focus is mostly on audio)

ASR and TTS: Defining an ASR model

We will use HuggingFace to get ASR and TTS models. Let's start with ASR

```python
from transformers import AutoModelForSpeechSeq2Seq, AutoProcessor, pipeline


def init_asr_model(config):
    torch_dtype = torch.float32 if config.device == "cpu" else torch.float16

    model = AutoModelForSpeechSeq2Seq.from_pretrained(
        config.asr.model_id,
        torch_dtype=torch_dtype,
        low_cpu_mem_usage=True,
        use_safetensors=True,
    )
    model.to(config.device)

    processor = AutoProcessor.from_pretrained(config.asr.model_id)

    asr_pipeline = pipeline(
        "automatic-speech-recognition",
        model=model,
        tokenizer=processor.tokenizer,
        feature_extractor=processor.feature_extractor,
        torch_dtype=torch_dtype,
        device=config.device,
    )
    return asr_pipeline

```

Model Id is the model name from the HuggingFace Hub. For example, *openai/whisper-tiny.en*

ASR and TTS: Running the ASR model

HuggingFace allows us to use pre-defined pipeline to simplify inference

```
28
29  def run_asr_model(asr_pipeline, audio):
30      # pipeline expects numpy array of shape (T)
31      # so we convert and take 0-th channel
32      text_output = asr_pipeline(audio.numpy()[0])["text"]
33
34      return text_output
35
```

ASR and TTS: Defining a TTS model

For the ASR, we used base classes. We can use the direct ones, as we show for TTS

```python
from transformers import (
    FastSpeech2ConformerHifiGan,
    FastSpeech2ConformerModel,
    FastSpeech2ConformerTokenizer,
)


def init_tts_model(config):
    tokenizer = FastSpeech2ConformerTokenizer.from_pretrained(
        "espnet/fastspeech2_conformer"
    )
    spec_generator = FastSpeech2ConformerModel.from_pretrained(
        "espnet/fastspeech2_conformer"
    )
    vocoder = FastSpeech2ConformerHifiGan.from_pretrained(
        "espnet/fastspeech2_conformer_hifigan"
    )

    return tokenizer, spec_generator, vocoder

```

Remember that TTS consists of two parts:

1. Creating MelSpec
2. Converting it to audio

Different models can be used for 1 and 2.

ASR and TTS: Running the TTS model

In contrast to ASR, here we define the pipeline ourselves:

```
41        # Generate audio
42        inputs = tokenizer(query, return_tensors="pt")
43        input_ids = inputs["input_ids"]
44        output_dict = spec_generator(input_ids, return_dict=True)
45        spectrogram = output_dict["spectrogram"]
46
47        audio = vocoder(spectrogram)
```

LLM

For the LLM, running inference on your own machine is complicated (not impossible though). It is better to use APIs.

We could use
well-known ChatGPT,
but its API is not free

LLM

For the LLM, running inference on your own machine is complicated (not impossible though). It is better to use APIs.

We could use well-known ChatGPT, but its API is not free

groq

The same interface but has free APIs. For example, for Meta's llama-3

LLM: init client

When we created an API key, we can start a LLM client:

```python
from groq import Groq

ROOT_PATH = Path(__file__).absolute().resolve().parent.parent


def init_llm():
    api_path = ROOT_PATH / ".env"
    with api_path.open("r") as f:
        api_key = f.readline().strip()
    client = Groq(api_key=api_key)

    history = [
        {
            "role": "system",
            "content": "You are an intelligent ai voice assistant named Sheila",
        }
    ]

    return client, history
```

Note that we need to handle history ourselves. We also start with a system message

LLM: run inference

For the inference, we just add the new query and delete too old elements from the history

```python
def run_llm(text, history, client, config):
    max_tokens = config.llm.max_tokens
    max_history = config.llm.max_history

    if len(history) >= max_history + 1:
        history = [history[0]] + history[-(max_history - 1) :]

    history.append({"role": "user", "content": text})

    chat_completion = client.chat.completions.create(
        messages=history,
        model="llama3-8b-8192",
        max_tokens=max_tokens,
    )

    predicted_text = chat_completion.choices[0].message.content

    history.append({"role": "assistant", "content": predicted_text})

    return predicted_text, history
```

Problem: Response may be looooooong

What is the LLM response is too long?

1. We cannot run TTS anymore (memory limitations)

2. Even if we can, generation of audio will take a lot of time. The user will wait in silence for a long period of time (it is bad)

Solution: split the response into several text partitions and generate audio for them individually

Solution

Split the text

```
26  def run_tts_model(text, tokenizer, spec_generator, vocoder, config):
27      max_words_per_query = config.tts.max_words_per_query
28
29      text_partitions = text.split()
30      returned_partitions = 0
31      number_of_partitions = len(text_partitions) // max_words_per_query
32      for i in range(number_of_partitions + 1):
33          if i * max_words_per_query == len(text_partitions):
34              break
35          query = text_partitions[i * max_words_per_query : (i + 1) * max_words_per_query]
36          returned_partitions += len(query)
37
38          query = " ".join(query)
39          print(f"{returned_partitions}/{len(text_partitions)}: {query}")
40
```

Solution

Process part of the text individually and yield it:

```
41        # Generate audio
42        inputs = tokenizer(query, return_tensors="pt")
43        input_ids = inputs["input_ids"]
44        output_dict = spec_generator(input_ids, return_dict=True)
45        spectrogram = output_dict["spectrogram"]
46
47        audio = vocoder(spectrogram)
48        yield audio
49
```

Writing to the loudspeaker

We could also create a separate process, but for the simplicity we are doing it in the main process.

```
99
100     stream_writer = StreamWriter(
101         dst=config.stream_writer.dst, format=config.stream_writer.format
102     )
103     stream_writer.add_audio_stream(
104         sample_rate=config.stream_writer.sample_rate, num_channels=1
105     )
106     stream_writer.open()
```

Writing to the loudspeaker

For each text partition, we get generated audio and play it chunk-by-chunk

```
118                # send it to ASR, LLM and TTS
119                user_full = torch.cat(user_chunks, dim=-1)
120                user_audio_output_generator, history = full_handler(
121                    all_models, history, user_full, config
122                )
123
124                start_time = time.perf_counter()
125                total_num_frames = 0
126                for user_audio_output in user_audio_output_generator:
127                    user_audio_output = user_audio_output[0].unsqueeze(-1)
128                    num_frames = user_audio_output.shape[0]
129                    total_num_frames += num_frames
130                    for i in range(0, num_frames, config.stream_writer.chunk_size):
131                        stream_writer.write_audio_chunk(
132                            0, user_audio_output[i : i + config.stream_writer.chunk_size]
133                        )
134                audio_time = total_num_frames / config.stream_writer.sample_rate
135                end_time = time.perf_counter()
136
```

Problem with writing

The process sends the chunks to the stream, but it does not know whether the loudspeakers have ended playing or not.

Solution: calculate the time of the audio response and make the main process sleep for these number of seconds.

```
134             audio_time = total_num_frames / config.stream_writer.sample_rate
135             end_time = time.perf_counter()
136
137             diff_time = audio_time - (end_time - start_time)
138             if diff_time > 0:
139                 print(audio_time, diff_time)
140                 time.sleep(1 + diff_time)  # to avoid ASR reading an output
```

Demo Time